

---

# Conceptualizing Surprise with the Jensen–Shannon distance\*

A Bayesian Information-Theoretic approach for the Social Sciences

Emil Niclas Meyer-Hansen<sup>†</sup>

14 January 2026

## Abstract

Generally in the social sciences, results are informally deemed *surprising* if their associated *p*-value is sufficiently small. This implicit interpretation, however, is both conceptually and mathematically inappropriate, and such misuse of the *p*-value can lead to erroneous conclusions about the novelty of results. To solve that issue, this paper builds on Bayesian inference, Information theory, and recommendations made to related fields, to argue for the adoption of a more appropriate conceptualization of surprise as the *relative entropy between prior and posterior knowledge*. This formal conceptualization enables researchers to appropriately measure surprise as the Jensen-Shannon distance, for which the paper provides easily implementable software and a demonstration of its use in relation to empirical data.©<sup>1</sup>

**Keywords:** *Surprise; Novelty; Jensen-Shannon distance; JS distance; Jensen-Shannon divergence; JS divergence; Relative entropy; Distance; Divergence; Dissimilarity; p-value; S-value; Information theory; Bayesian inference*

---

\***Version:** 2026-01-14-15-45. **Citation:** Meyer-Hansen, E. N. (2026): ‘Conceptualizing Surprise with the Jensen–Shannon distance: A Bayesian Information-Theoretic approach for the Social Sciences’, *Open Science Framework*, Working paper (v2026-01-14-15-45). DOI: 10.17605/OSF.IO/GQ6C8

<sup>†</sup>**Author:** Independent Researcher, Aarhus, Denmark. Holds a Master of Science (MSc.) and Bachelor of Science (BSc.) in Political Science from Aarhus University, Denmark (email: emil098meyerhansen@gmail.com).

<sup>1</sup>**License:** Except where otherwise indicated, all contents of this document and associated files are licensed under the *Creative Commons Attribution 4.0 International License* (CC BY 4.0). All software, including but *not* necessarily limited to, source code, executable code, code snippets, code chunks, algorithms, and/or scripts, attributable to this document and/or its associated files are expressly excluded from the foregoing license, and unless otherwise indicated, are instead licensed under the *GNU General Public License, version 3* (GPL-3.0). By engaging with this document and/or any associated files, which include, but are *not* necessarily limited to, downloading, using, viewing, and/or distributing any of them, in parts or whole, you agree to comply with the applicable license terms for the respective content types.

## 1 Introduction

For a research paper to be deemed worthy of publication in a peer-reviewed scientific journal, besides adhering to standards on rigor and replicability (Christensen et al., 2019; Shaughnessy et al., 2015; Toshkov, 2016), its results should ideally also be *surprising* (or novel) when compared to expectations based on intuition or theory (French and Murphy, 2023; Morgan, 2005; Wang et al., 2024). This preference for surprise/novelty could be attributed to the belief that scientific studies should serve to produce knowledge relevant for solving societal problems (Buunk and van Vugt, 2013; Lugg, 1979), which implies a limitation of current problem-solving capabilities and thus the need for a study to acquire *new* knowledge. Similarly, from the perspective of editors, if the goal is to increase the readership of a journal, novelty may be preferred, since surprising results are more likely to attract attention and engage readers (cf. Itti and Baldi, 2009). To avoid their research collecting dust in the file-drawer (cf. Rosenthal, 1979), the surprise of results thus becomes an important factor for researchers to assess and emphasize in order to increase the chances of publication.

Recognizing this preference, the scientific community has tried to standardize and quantify novelty to make it less subjective (e.g., Modirshanechi et al., 2021; Wang et al., 2024), but despite their best efforts, there is currently no broad consensus, leaving different fields to rely on different solutions (Wang et al., 2024). Within the social sciences, this has lead to the development of informal practices for assessing novelty, with particularly the (in)famous *p*-value currently serving as a popular heuristic. Here, an informal assessment of surprise involves checking whether the *p*-value is smaller than some threshold (e.g.,  $p \leq 0.05$ ), and if so, the results are implicitly considered to be 'surprising'. This means that the *p*-value is effectively treated as a measure of surprise within the social sciences.

However, using the *p*-value as a measure of novelty can be shown to be inappropriate given its conceptualization and mathematical properties. Specifically, it can be made evident that almost any result becomes 'surprising' given a sufficiently large sample size, irrespective of how derivative the results actually are compared to the expectations of intuition or existing theory. Misleading conclusions about the novelty of results are thus abound to incur from this informal practice, and it therefore constitutes an issue for the social sciences that may lead to incorrect decisions in relation to publication.

This paper argues for the need to adopt a formal (re)conceptualization of surprise within the social sciences to solve this issue. Building on Bayesian inference and Information theory, this involves conceptualizing *knowledge* as *evidence-based beliefs*, where the expected results given existing theory can be characterized as *prior* knowledge, while the actual results given the data can be characterized as *posterior* knowledge. This enables a theoretically-coherent, formal conceptualization of surprise as *the relative entropy between prior and posterior knowledge*. As demonstrated in relation to empirical data, surprise can then be appropriately measured as the *Jensen-Shannon distance*, which possess numerous advantages suited for the social sciences. By arguing for this conceptualization, the paper thus adds to the growing literature that calls for formally conceptualizing surprise in accordance with Bayesian inference and Information theory (e.g., Bencomo and Belaggoun, 2014; Burkhart et al., 2025; Burnham and Anderson, 2001; Itti and Baldi, 2009; Mello et al., 2025; Seehars et al., 2014).

To achieve this solution, the paper is structured as the following: Following this introduction (Section 1), the paper considers the current problem within the social sciences by reviewing the informal conceptualization of

surprise and the *p*-value to argue for its (mis)use as an implicit indicator of surprise within the social science literature (Section 2). The paper then builds towards a solution by presenting concepts central to Bayesian inference and Information theory and incorporating them into a formal conceptualization of surprise as the relative entropy between prior and posterior knowledge, measurable as the Jensen-Shannon distance (Section 3). A practical implementation of the Jensen-Shannon distance for popular statistical software is presented, before a demonstration of its ease of use in relation to empirical data is provided (Section 4). The uses and limitations of this solution are discussed (Section 5) before the paper is concluded (Section 6).

## 2 Problem

To build towards a formal conceptualization of surprise as the relative entropy between prior and posterior knowledge, this section first covers the current, informal definition of surprise within the social sciences (Section 2.1), before reviewing the conceptualization and mathematical properties of the *p*-value that make it inappropriate as a measure of surprise (Section 2.2).

### 2.1 Surprise, informally

In layman's terms, *surprise* can generally be defined as a reaction 'caused by something unexpected happening' (Cambridge Dictionary, 2025). While this definition is consistent with researchers' application of the term (cf. Ivanova and Vaidyanathan, 2024), including within the social sciences (e.g., Alford et al., 2011: 362; Goren and Chapp, 2017: 124; Petersen and Laustsen, 2019: 26), it can be argued to be too vague to constitute a formal scientific conceptualization. However, usage of this term by researchers can nevertheless be analyzed as a starting point to understand why the *p*-value is misused as a measure of surprise and what a proper conceptualization within the social sciences would need to consider.

That surprise is a reaction to the unexpected can within the social sciences be understood to reflect a *discrepancy between results and predictions made about those results based on theory*. Here, predictions are generally formalized as *hypotheses* (Shaughnessy et al., 2015; Toshkov, 2016), whose theoretical basis may range from weak (e.g., intuition-based) to strong (e.g., well-established theory). That such discrepancy is (informally) recognized as being synonymous with surprise is made evident by differences between expectations and results almost ubiquitously being referred to as surprising and worthy of discussion. For example, Cochrane et al. (1979) studied the political values of individuals in British neighborhoods with door-to-door interviews. Based on theory by Rokeach (1973), they expected individuals with different political ideologies (e.g., liberals, communists, fascists) to differ in their prioritization of freedom. Instead, they found that individuals across all ideologies prioritized freedom similarly, and this unexpected discrepancy between the predictions by Rokeach (1973) and the results by Cochrane et al. (1979) were discussed and later explicitly referred to as 'surprising' by other researchers (e.g., Sterling et al., 2019: 2). Surprise can thus informally be characterized as the discrepancy between results and theory, and a formal conceptualization of it would thus need to reflect this property.

Another related characteristic of surprise that may be recognized from its informal usage is that it is entirely *relative* to the predictions of the theory to which results are being compared. This is made evident by researchers recognizing that results may be surprising when evaluated against one theory, but that they can also be entirely unsurprising when compared to another. For example, Osmundsen et al. (2021) investigate the behavior of Ameri-

can individuals on the social media *X* (then known as *Twitter*). In contrast to the expectations of *ignorance theory*, where fake news sharing would expectedly be attributable to individuals being unable to discern between true and fake information (Osmundsen et al., 2021: 1001; based on Altay et al., 2020; Pennycook and Rand, 2019; Pennycook et al., 2020), they instead found that it was better explained as based in hostility toward political opponents in line with the *negative partisanship* literature (Osmundsen et al., 2021: 1012; see Abramowitz and Webster, 2018). While this result could thus be considered surprising when compared to one theory, it was entirely unsurprising when compared to expectations of another. This means that an appropriate conceptualization of surprise needs to account for it being *relative* to the different theories that results are being compared against.

A final characteristic of surprise gained from analyzing its popular use is that it is *dynamic* and evolves as science progresses. This is made evident by researchers using surprising results to argue for revisions, if not discardment, of theory, with the aim being to bring expectations in greater alignment with results and thereby reduce future surprise. For example, in the discussion of the previously noted study by Osmundsen et al. (2021), they use their results to argue for revising future expectations about fake news sharing, stating that 'If people care primarily about a story's ability to hurt political enemies, we should not be surprised if fact checking fails to reduce sharing.' (Osmundsen et al., 2021: 1013) Such revisions fit the self-correcting ideal of science, whereby theories are revised or discarded in favor of theories with a higher explanatory power (i.e., *verisimilitude*, Miller, 1974; Popper, 1959; see also Dellsén, 2024; Lakatos, 1970). In the event that similar results were to appear in future studies, they should accordingly incur *less* surprise. This example also corroborates the notion that, whether phrased in terms of 'weak' explanatory power, 'bad' fit, or a similar term, discrepancy between results and predictions are indicative of the *epistemic limitations* of the theory underlying those predictions<sup>2</sup> (French and Murphy, 2023: 1448-1449). As such, an appropriate conceptualization of surprise would enable researchers to motivate revisions or discardment of theory, and that the better the predictions of a theory, the less surprising the results should appear.

Based on these initial considerations, it is apparent that surprise can currently be considered vaguely defined within the social sciences. Nonetheless, based on an analysis of its common usage, it has become clear that it can be characterized as the discrepancy between results and predictions derived from theory, with surprise being relative to, and co-evolving with, theory, with higher degrees of surprise being indicative of its limitations. These insights will prove important for any formal conceptualization of surprise within the social sciences, since a proper measure of surprise will need to account for these characteristics to be suitable for that discipline. While some may argue that a vague conceptualization of surprise is sufficient for the social sciences, the pitfalls of this informality is demonstrated in the following section by considering the misuse of the *p*-value as an implicit measure of surprise.

## 2.2 *p*-value

In the classical statistical framework, also known as the *Frequentist* approach (for an introductory text, see, e.g., Agresti, 2018), the *probability value* (i.e., *p*-value) denotes the *probability of at least as extreme an estimated parameter value from the sample given that the null-hypothesis about the population parameter is true* (Clayton, 2021: 291). This dense definition presupposes some conceptual understanding, which can be briefly explained.

The *p*-value is conceptually part of the *Null-Hypothesis Significance Testing* procedure (NHST; Cohen, 1994;

<sup>2</sup>Indeed, at least in a deterministic world, a researcher making predictions using a theory with perfect explanatory power will never be surprised by the results. As a note of reference, it may thus be stated that Laplace's (1951) demon is incapable of being surprised (for arguments against a deterministic world, see, e.g., Masi, 2023).

Field, 2018: 72-82; Kruschke, 2014: 297-334), which involves the researcher formulating a *null-hypothesis* (i.e.,  $H_0$ ) about a parameter in a population of interest (i.e.,  $\theta$ ), which typically takes the form of a *nil hypothesis* by stating that the parameter is exactly zero (i.e.,  $H_0 : \theta = 0$ ; Cohen, 1994: 999-1000). This null-hypothesis contrasts an *alternative hypothesis* (i.e.,  $H_a$ ) about the population parameter that the researcher has derived from theory, though it generally is just a non-nil hypothesis (e.g.,  $H_a : \theta \neq 0$ ).

The purpose of NHST is then to build confidence for  $H_a$  through attempts at falsifying  $H_0$ . This is accomplished by collecting a sample that is sufficiently large and representative of the population of interest, before calculating the parameter from this sample using a mathematical formula known as an *estimator* (Stock and Watson, 2019: 105; Wooldridge, 2019: 715-721), which provides an estimate of the parameter in the population based on the sample. The estimated parameter from that sample (i.e.,  $\hat{\theta}_s$ ) is then compared to the *sampling distribution* of its estimator (i.e.,  $\hat{\theta}$ ; Field, 2018: 61-64), which is a distribution of hypothetical estimates derived from an infinite number of alternative samples taken from that population if  $H_0$  is true. Since Frequentists define probability as *the long-run relative frequency* (Agresti, 2018: 79), the *p-value* is calculated as the area under the curve in this sampling distribution of  $\hat{\theta}$  that is at least as extreme as  $\hat{\theta}_s$ . If the *p-value* is lower than some threshold (i.e.,  $\alpha$ ) specified in advance (e.g.,  $\alpha = 0.05$ , with  $p \leq \alpha$ ) then  $H_0$  is rejected, which would be consistent with  $H_a$  being true, and the estimated parameter is referred to as 'statistically significant'. By contrast, if the *p-value* is higher than  $\alpha$ , then  $H_0$  fails to be rejected, which is consistent with  $H_0$  being true, and the estimated parameter is referred to as 'statistically insignificant'.

From this conceptualization of the *p-value*, it is evident that the intended purpose of this statistic is that of *statistical inference*, that is, inferring results from sample to population (Agresti, 2018: 16-17). Nowhere in its conceptualization is a consideration for expressing the degree of discrepancy between results and predictions based in theory (i.e., either a small or large *p-value* could be consistent with any degree of actual surprise). While the *p-value* is relative to a prediction based in the null-hypothesis, the null-hypothesis generally does not reflect serious predictions based in theory (i.e., what social theory would be considered substantive if its sole prediction involved a parameter-value being non-zero?), and the null-hypothesis generally fails to evolve with a theory by generally remaining a nil hypothesis no matter the context, making the *p-value* too static to properly reflect the dynamic nature of surprise.

These issues can be further recognized with a simple example. Suppose one conducts an experiment on the entire population of interest, where half of individuals are exposed either to a placebo condition or some treatment, which given one theory is expected to change these individuals, while another theory expects it not to change anything. Even if the treatment is found to induce a change or not, the *p-value* is entirely irrelevant to the purposes of this experiment, since there is no need for statistical inference (because the sample is the entire population).<sup>3</sup> Since an assessment of how surprising the results are remains relevant in spite of the irrelevance of the *p-value*, this means that an assessment of surprise must be separate from it.

These conceptual limitations can be further realized by considering the mathematical formula underlying the *p-value*. While it can be calculated in numerous ways, the most widely used version of the *p-value* involves an assumption that the sampling distribution of the estimator used to calculate the parameter of interest ( $\hat{\theta}$ ) is

---

<sup>3</sup>While one may argue that calculating the *p-value* in this example could help express uncertainty, it should be emphasized that the only uncertainty that can be expressed by the *p-value* is that of *sampling uncertainty*, which is only relevant for statistical inference and therefore irrelevant in the example.

*Gaussian* (i.e., normal; Wooldridge, 2019: 723-724), which under a nil hypothesis, means that the distribution can be characterized with the following zero-mean normal probability distribution ( $\mathcal{N}$ ; Agresti, 2018: 84-91):

$$\hat{\theta} \sim \mathcal{N}(0, \sigma_{\hat{\theta}}) = \frac{1}{\sqrt{2\pi\sigma_{\hat{\theta}}^2}} e^{-\frac{x^2}{2\sigma_{\hat{\theta}}^2}} \quad (1)$$

where  $\sigma_{\hat{\theta}}$  refers to the dispersion of this sampling distribution of  $\hat{\theta}$  (i.e., the *standard error*; Field, 2018: 61-64);  $\pi = 3.141 \dots$  is the ratio between the circumference and diameter of a circle; and  $e = \sum_{n=0}^{\infty} \frac{1}{n!} = 2.718 \dots$  is the base of the exponential function.

If  $H_0 : \theta = 0$  and  $H_a : \theta \neq 0$ , the direction by which  $\hat{\theta}_s$  deviates from 0 is irrelevant, and the *p-value* is accordingly referred to as *two-tailed* (Field, 2018: 79-81), which given the symmetry of the distribution in (1) is given by the definite integral (Stock and Watson, 2019: 111):

$$p \equiv 2 \int_{|\hat{\theta}_s|}^{\infty} \mathcal{N}(0, \sigma_{\hat{\theta}}) dx \quad (2)$$

where  $\int$  denotes the integration operator and  $|\hat{\theta}_s|$  is the absolute value of  $\hat{\theta}_s$ .

Since an application of formula (2) requires calculating  $\sigma_{\hat{\theta}}$  using infinitely many samples from a population in the possibly hypothetical scenario where  $H_0$  is true,  $\sigma_{\hat{\theta}}$  is generally estimated. For example, if the parameter of interest is the population mean (i.e.,  $\theta = \mu$ ) with  $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$  being its estimator (Field, 2018: 27), where  $x_i$  denotes the *i*th value from an *n*-length vector constituting a random sample of independent and identically distributed (*iid*) values drawn from a normally distributed population of interest; and the population variance is estimated as  $\hat{\sigma}^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \hat{\mu})^2$  (Field, 2018: 30), then the estimator for the dispersion of the sampling distribution for the estimated mean is (Stock and Watson, 2019: 113, equation 3.9):  $\hat{\sigma}_{\hat{\mu}} = \frac{\hat{\sigma}}{\sqrt{n}}$ .

Now, since  $\hat{\sigma}^2$  is an *unbiased* and *consistent* estimator of  $\sigma^2$  (Stock and Watson, 2019: 112; Wooldridge, 2019: 716-723, 766), they converge as the sample size (*n*) increases towards infinity, while  $\sqrt{n}$  increases without bound, and a mathematical analysis of  $\hat{\sigma}_{\hat{\mu}}$  thus makes it evident that the estimated dispersion of the sampling distribution approaches zero as *n* increases towards infinity:

$$\lim_{n \rightarrow \infty} \frac{\hat{\sigma}}{\sqrt{n}} = 0 \quad (3)$$

The implications of equation (3) pose a detrimental issue for the ability of the *p-value* to reflect surprise. This can be made more clear by first converting the estimated parameter into a standardized format called a *test statistic* (Field, 2018: 78-79). In relation to an assumption of a normal sampling distribution, this test statistic is referred to as a *z-value*, which involves the rescaling formula (Field, 2018: 37):  $z = \frac{\hat{\mu} - \mu_{H_0}}{\hat{\sigma}_{\hat{\mu}}}$ . If the difference between the population parameter and the population parameter under the null-hypothesis is non-zero (i.e.,  $\mu - \mu_{H_0} \neq 0$ ), this *z*-value can be recognized to approach infinity as  $\hat{\sigma}_{\hat{\mu}}$  decreases towards zero, since  $\hat{\mu} - \mu_{H_0}$  converges to  $\mu - \mu_{H_0}$  as *n* increases towards infinity due to  $\hat{\mu}$  being unbiased, consistent, and otherwise independent of  $\hat{\sigma}_{\hat{\mu}}$ , while  $\mu_{H_0}$  is a constant (Stock and Watson, 2019: 104-107; Wooldridge, 2019: 716-723). Similar to (3), this can be expressed as:

$$\lim_{\hat{\sigma}_{\hat{\mu}} \rightarrow 0} \frac{\hat{\mu} - \mu_{H_0}}{\hat{\sigma}_{\hat{\mu}}} = \infty \quad (4)$$

which has been established to occur in (3) when  $n \rightarrow \infty$ . Due to the  $z$ -standardization, the normal distribution in (2) has become a standard normal, and substituting  $\hat{\theta}_s$  for  $z$ , it can finally be realized that the limit of the  $p$ -value as  $z$  increases towards infinity is zero:

$$\lim_{z \rightarrow \infty} 2 \int_{|z|}^{\infty} \mathcal{N}(0, 1) dx = 0 \quad (5)$$

as demonstrated to occur in (4) when  $\hat{\sigma}_{\hat{\mu}} \rightarrow 0$ , which happens when  $\mu - \mu_{H_0} \neq 0$  and  $n \rightarrow \infty$  as shown in (3).

The implication of these limits is that for any instance where the null-hypothesis is not *exactly* true — which likely applies to most of the phenomena studied in the social sciences<sup>4</sup> — that anything becomes 'statistically significant' with a sufficiently large sample size. In the event that one equates a low  $p$ -value with surprise, this unfortunately means that *almost anything becomes surprising with a large enough sample size*, irrespective of how expected the results actually are under the given theory.<sup>5</sup> Using the  $p$ -value as a measure of surprise is thus conceptually flawed and runs counter to the characteristics of actual surprise recognized by social science researchers as previously in section (2.1).

While one may argue that the informal practice of using the  $p$ -value as a measure of surprise is too negligible, it may be counterargued that such misuse is best addressed as early as possible, so as to prevent it from spreading so wide that it becomes formal. Such a development may seem improbable, but this happened recently within the field of epidemiology. Here, Cole et al. (2021) introduced the '*S*-value' as a formal measure of surprise. Researchers, however, could quickly realize that this '*S*'-value unfortunately had a one-to-one correspondence with the  $p$ -value through the formula:  $S \equiv -\log_2(p)$  (Cole et al., 2021: 192). Similar to criticisms made against other measures derived solely from the  $p$ -value<sup>6</sup>, this means that the *S*-value fails to incorporate, nor communicate, any new information besides that already expressed by the  $p$ -value. The *S*-value is thus as susceptible to the exact same issues identified here for the  $p$ -value in its ability to express surprise<sup>7</sup>, making it just as *inappropriate* for such use.

Now, while this paper serves to solve this misuse of the  $p$ -value, the aim is *not* to portray researchers negatively for partaking in this inappropriate practice. Indeed, the  $p$ -value is infamous for having been misinterpreted in numerous ways (Clayton, 2021: 291; Gigerenzer, 2018: 206), including as an index of scientific significance,

---

<sup>4</sup>That the null-hypothesis is generally false within the social sciences can be realized by considering that hypotheses generally concern point values (e.g.,  $H_0 : \theta = 0$ ) despite most studied phenomena being continuous, which involve types of distributions where the hypothesis has an *a priori* zero probability of being true (cf. Cohen, 1994: 1000; Tukey, 1991: 100). In psychology, the notion that the nil hypothesis is always false is referred to as Meehl's (1978) conjecture (Gigerenzer, 2004: 601).

<sup>5</sup>The issue cannot be addressed by specifying the null-hypothesis to fit the prediction of a given theory, since the most negligible difference between the result and the prediction can produce a vanishingly small  $p$ -value given a sufficiently large sample.

<sup>6</sup>Somewhat unsurprisingly, this is not the first time that the  $p$ -value has been inappropriately repurposed to express phenomena outside its intended purpose. As an infamous example, Killeen (2005a) similarly sought to address the fact that the  $p$ -value cannot express the probability that the results of a study will replicate in another study (cf. Gigerenzer, 2018). The mathematical formula underlying the proposed solution was widely criticized (e.g., Iverson et al., 2009; Macdonald, 2005; Maraun and Gabriel, 2010; Serlin, 2010; for the response, see, Killeen, 2005b), since it was evident that the formula had a one-to-one correspondence with the  $p$ -value, and it thus failed to incorporate any information about the probability of replication not already communicated by the  $p$ -value. Instead, as similarly argued here in relation to surprise, it was argued that a (re)conceptualization of replicability was better founded in Bayesian inference (e.g., Macdonald, 2005).

<sup>7</sup>This can be made evident by considering that  $\lim_{p \rightarrow 0} -\log_2(p) = \infty$ , which was established by equations (3) through (5) to occur as  $n \rightarrow \infty$ , which means that, based on the logic of the *S*-value, almost anything will inevitably become surprising with a sufficiently large sample size, no matter how expected the results were in advance. Despite these issues, other researchers have worryingly advocated for the adoption of the *S*-value (e.g., Greenland et al., 2022; Rafi and Greenland, 2020), further revealing the need for a proper formal (re)conceptualization of surprise.

reliability, replicability, effectsize, the probability of  $H_0$ ,  $H_a$ , or data having been generated by chance, resulting in various criticisms and cautions against its (mis)use (American Statistical Association, 2016; Clayton, 2021; Field, 2018: 97-110; Gelman and Stern, 2006; Kruschke, 2014: 297-329; McShane et al., 2019). Instead, as similarly achieved by other researchers identifying questionable practices within the scientific community (e.g., Bishop, 2006; Flake and Fried, 2020; Gelman, 2024; Ioannidis et al., 2017; Lenz and Sahn, 2021; Ritchie, 2020; Simmons et al., 2011), the aim is to improve the social scientific practice by informing researchers of inappropriate practices and to propose possible solutions.

### 3 Solution

While the previous section sought to describe the current, informal conceptualization of surprise and show how the (mis)use of  $p$ -value as a measure of surprise is inappropriate and constitutes an issue for the social sciences, this section serves to provide a solution to that problem with a formal (re)conceptualization of surprise. To that end, central concepts of Bayesian inference (Section 3.1) and Information theory (Section 3.2) are presented and used to build that conceptualization and identify an appropriate measure for it.

#### 3.1 Bayesian inference

The *Bayesian statistical framework*, also known as Bayesian inference, is an approach to statistical inference that differs from its Frequentist counterpart by defining probability in relation to (inter-)subjective beliefs (for an introduction, see, e.g., McElreath, 2019). Understanding this framework presupposes some familiarity with *Bayes' theorem* (Bayes and Price, 1763; Laplace, 1951), which can be expressed in relation to the probability of two discrete events as (McElreath, 2019: 49):

$$P(A|B) = \frac{P(A)P(B|A)}{P(B)} \quad (6)$$

where  $P(A)$  denotes the probability of the event  $A$ , which is referred to as the *prior* probability;  $P(B)$  denotes the probability of event  $B$ , sometimes called the *marginal likelihood*;  $P(B|A)$  is the probability of event  $B$  given that event  $A$  has occurred, denoted the *likelihood*; and  $P(A|B)$  is the probability of event  $A$  given that event  $B$  has occurred, which is referred to as the *posterior* probability.

The components of (6) are generally reconceptualized for the purposes of statistical inference (Gelman et al., 2014; Kruschke, 2014; Levy and Mislevy, 2020), so that the parameter of interest ( $\theta$ ) substitutes  $A$ , while the collected data ( $\mathcal{D}$ ) substitutes  $B$ , which means that  $P(\theta)$  denotes the (prior) probability of  $\theta$ ,  $P(\mathcal{D}|\theta)$  the probability of the data given the parameter, and  $P(\theta|\mathcal{D})$  the (posterior) probability of  $\theta$  given  $\mathcal{D}$ , with this latter component being of primary interest in Bayesian inference. For practical purposes  $P(\mathcal{D})$  can often be ignored due to relying on *Markov-Chain Monte Carlo methods* (MCMC; Brooks et al., 2011), which leaves  $P(\theta)$  and  $P(\mathcal{D}|\theta)$  relevant for calculating  $P(\theta|\mathcal{D})$ . While those familiar with *likelihood-based inference* (King, 1998) will know that  $P(\mathcal{D}|\theta)$  is often easy to derive from the statistical model used in the analysis,  $P(\theta)$  is by comparison rarely known. To solve that issue, this prior is specified as a probability distribution using *a priori* knowledge about  $\theta$ . The discretion involved in specifying this prior distribution is where Bayesian inference becomes (inter-)subjective, though results generally approximate Frequentist results with a sufficiently large sample size (Cover and Thomas, 2006:

388; Hastie et al., 2017: 272; King, 1998: 30; Kruschke, 2014: 113), which means that Bayesian inference could be considered a generalization of the arguably less subjective<sup>8</sup> likelihood-based approach (cf. King, 1998: 28). As such, the difference between the two approaches is primarily in the definition, and thus interpretation, of probability.

This conceptualization enables Bayesian inference to answer questions about the probability of parameters given the data,  $P(\theta|\mathcal{D})$ , which is impossible in the Frequentist framework, since that approach solely relies on the likelihood and fails to consider prior information in relation to  $\theta$  (Kruschke, 2014; Wagenmakers et al., 2010). It also means that results from a Bayesian analysis are in the form of probability distributions, which conceptually express how credible different estimated values of  $\theta$  are given the data. While these distributions could be summarized with point estimates (e.g., mean, dispersion) or intervals (e.g., 95% credibility interval) for practical purposes, as done in the Frequentist approach, this is theoretically unnecessary, since the posterior distribution comprehensively details one's knowledge about  $\theta$  based on the prior knowledge and data (McElreath, 2019: 58).

Now, the conceptualization of Bayes theorem used for statistical inference can be recognized to naturally capture the informally recognized characteristics of surprise identified in section (2.1). With surprise being the discrepancy between results and expectations/predictions based in theory, it can be equivalently understood as a discrepancy between the state of knowledge *prior* to seeing the results and the state of knowledge *posterior* to seeing the results. When theory is formalized and well-established among peers within the social sciences, the prior thus characterizes the state of knowledge as a relatively intersubjective belief about a phenomenon, describing what the researcher expects the results to be. As established, since the prior distribution is conceptually a comprehensive statement about this former state of knowledge, it can be appropriately used to characterize the expectations based in theory.

### 3.1.1 The prior reflects current knowledge

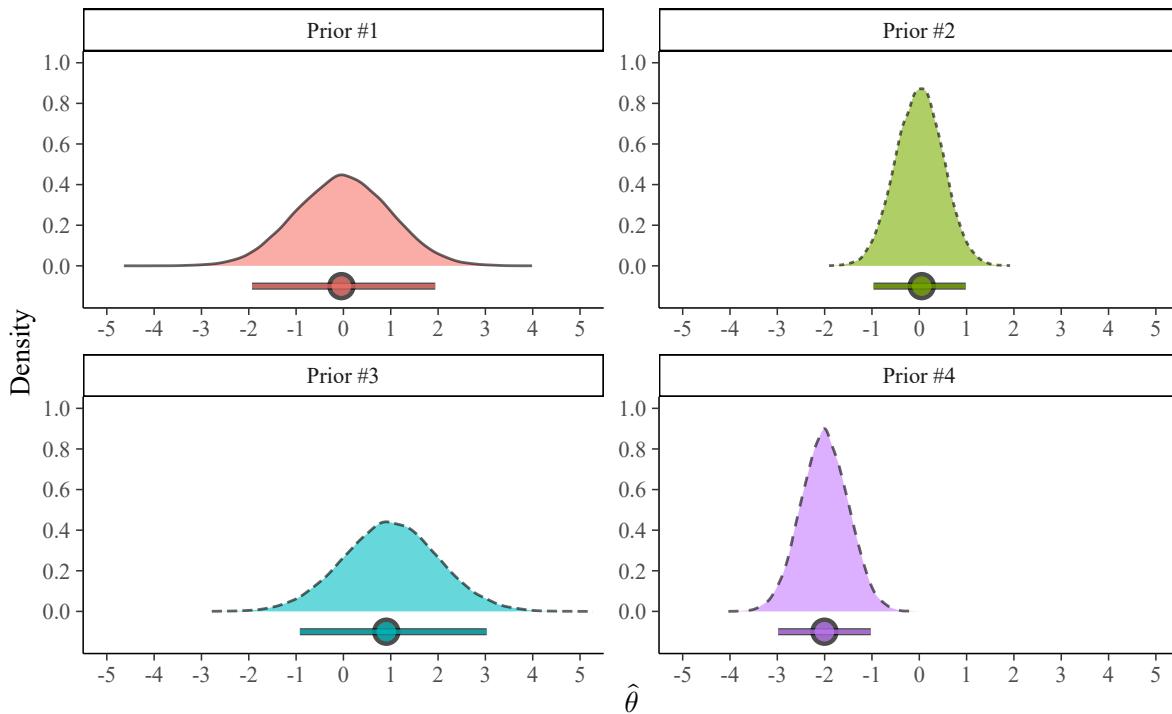
Characterizing the state of knowledge in a theory as a prior distribution can be difficult, especially if it only makes vague predictions about the parameters relevant to that theory. To therefore help build intuition around them, four examples of different specifications are provided here.

At no loss of generality, suppose  $X$  and  $Y$  are two standard normally distributed variables that are theorized to be linearly related in some population of interest:  $Y_i \sim \mathcal{N}(\theta X_i, \sigma)$  for all  $i$  in  $1, \dots, N$ , where one wishes to estimate the coefficient  $\theta$  using a representative and sufficiently large sample. To enable Bayesian inference, it is necessary to specify its prior,  $P(\theta)$ , which can be achieved in numerous ways depending on one's prior information about it. As shown visually in figure (1), the first prior (Prior #1) is specified as a standard normal distribution:  $\mathcal{N}(0, 1)$ . Such a prior can be taken to reflect a diffuse state of knowledge, where there is complete theoretical ignorance about the direction with which a parameter may influence a relationship, but its relatively substantial width (i.e., size of the dispersion) implies that the parameter could be plausibly expected to have some influence. This prior may be referred to as 'skeptical, yet persuadable', which is a type of weakly informative prior that is critical with respect to the sign of a parameter of interest, but which remain wide enough to be influenced by the results. The second prior (Prior #2) is similar to the first, but by comparison, it is much narrower, which implies

---

<sup>8</sup>Since all human processing of information and knowledge is necessarily (inter-)subjectively dependent (cf. Husserl, 1983; see also Kant, 1781), the Frequentist statistical framework, including its likelihood-based approach, is *not* 'objective' (i.e., independent of (inter-)subjective experience), despite claims to the contrary (e.g., Mayo, 1996: 89).

a stronger expectation that the parameter does not have an influence on the relationship. This contrasts the third prior (Prior #3), which has the expectation that the influence of the parameter on the relationship is positive, but its similar width to the first prior means it reflects a less certain expectation. The last prior (Prior #4) characterizes an opposite expectation to the third prior, since it expects the influence on the relationship to be negative, and its greater narrowness reflects a greater degree of certainty about this theoretical belief.



**Figure 1.** Example of Distributions

**NOTE:** Prior #1  $\sim \mathcal{N}(0, 1)$ ; Prior #2  $\sim \mathcal{N}(0, 0.5)$ ; Prior #3  $\sim \mathcal{N}(1, 1)$ ; Prior #4  $\sim \mathcal{N}(-2, 0.5)$ . Samples = 40,005. The geometrics below the distributions indicate the mode (circle) and the 95% HDCI (bar). Based on simulated data.

### 3.1.2 The posterior as updated knowledge

While the state of knowledge prior to the results could be difficult to precisely express as a probability distribution, there is no need for the researcher to specify the state of knowledge posterior to the results, since this resulting probability distribution is a function of the chosen prior distribution and the data. This means that the posterior distribution characterizes how the predictions have changed in light of the data. It can thus be considered an update state of knowledge, which is useful, since it is informative about how the expectations of that theory should change based on the data to enable better predictions in the future.

Since the posterior is the product of the prior and the data, this conceptualization helps reflect the aforementioned characteristics of surprise as being relative and dynamic in relation to the theories to which results are being compared. Since different prior specifications can reflect the varying predictions made across theories, the posterior could theoretically be compared to each to enable an assessment of how 'surprising' the results are given each theory. For example, let the posterior distribution resulting from an analysis be identical to the fourth example of the priors (Prior #4) previously shown in figure (1), meaning it is characterized as a normal probability distribution with a mean parameter of  $-2$  and a dispersion of  $0.5$ :  $\mathcal{N}(-2, 0.5)$ . Assume that these aforementioned priors

represent a set of predictions made from different theories, and that the data is large enough to make the posterior independent of these priors. Comparing the posterior to each prior then makes it clear that the discrepancies between them are not equivalent. When evaluated against the fourth prior (Prior #4), the results fit perfectly with the predictions made *a priori* about the parameter, which would be fully consistent with a complete lack of surprise based on how social science researchers use the term. While the difference between the posterior mean (i.e., -2) and the mean of priors one and two (i.e., 0) is the same, the results, however, would intuitively appear more 'surprising' when compared against the second prior, since that distribution is more narrow and thus expressed greater certainty about the parameter values as implied by its lower dispersion (i.e., 0.5 vs. 1). As such, while the third prior has the highest discrepancy with the posterior in terms of the mean, the dispersion of the distribution should also be considered when evaluating the degree of surprise under this Bayesian conceptualization, which will prove important when moving from a conceptualization of surprise towards an appropriate measurement of it.

Using these concepts from Bayesian inference, surprise has thus far been conceptualized as a discrepancy between states of knowledge, with these states specifically being expressible as prior and posterior probability distributions. The advantage of this approach has been demonstrated to arise from the conceptual properties of Bayesian inference that enable comprehensive statements to be made with relative ease about such states of knowledge. A shortcoming of this approach, however, is that Bayesian inference can be considered conceptually weak of terms of meaningfully quantifying the degree of discrepancy between states of knowledge, which is why the following section turns to the key concepts of Information theory that will build upon this Bayesian-based conceptualization to enable an appropriate measurement of surprise.

## 3.2 Information theory

Information theory is a relatively recent mathematical field that is primarily concerned with the efficient communication and compression of information (e.g., Hartley, 1928; Nyquist, 1924, 1928; Shannon, 1948, 1949; Wiener, 1961), but its concepts have proven useful across numerous disciplines, including Bayesian inference (e.g., Chaloner and Verdinelli, 1995; Jaynes, 1968). Being closely related to statistical mechanics (Jaynes, 1957a; 1957b), information is generally conceptualized within this discipline as the entropy of a random variable, which is the key concept in which an appropriate measure of surprise is founded. Since this and related concepts have so far seen little use within the social sciences, they are briefly presented here.

### 3.2.1 Entropy

In Information theory, *entropy* is 'a measure of the average uncertainty in the random variable' (Cover and Thomas, 2006: 5; Shannon, 1948), making it synonymous with randomness, unpredictability, and disorganization (Berrett et al., 2019: 288; Cover and Thomas, 2006: 135; Jaynes, 1957a: 622; Rényi, 1961: 547). It can be thought of as the information needed, on average, to describe a random variable (Berrett et al., 2019: 288; Cover and Thomas, 2006: 19), which means that the more random the variable, the higher its entropy.

For a finite discrete random variable  $X$ , the *Shannon-Boltzmann entropy*,  $H(X)$ , is defined by the mathematical formula (Back et al., 2018: 1; Cover and Thomas, 2006: 5, equation 1.1; Rényi, 1961: 547):

$$H(X) \equiv \sum_{x \in X} p(x) \log \frac{1}{p(x)} \quad (7)$$

$$= - \sum_{x \in X} p(x) \log p(x) \quad (8)$$

where  $\sum$  denotes the summation operator applied over the unique elements of  $X$ ;  $p(x) = P(X = x) \geq 0$  is the probability mass function for  $X$ ; and  $\log$  is the logarithmic function, where it is taken by convention that  $0 \log 0 \equiv 0$  (Cover and Thomas, 2006: 14). Since the base of the logarithmic is generally taken as 2, this means that the scale of entropy is in *binary units*<sup>9</sup> (i.e., *bits*, which are *not* to be confused with the related concept of binary digits; Cover and Thomas, 2006; Delgado-Bonal and Marshak, 2019; Shannon, 1948). Being measured in bits means that the entropy can be meaningfully interpreted as the expected number of binary questions (Yes/No) needed to know a value of  $X$  with absolute certainty (Cover and Thomas, 2006: 16).

With this conceptualization, the entropy can be stated to be a non-negative real number with the general upper bound being the logarithm of the number of unique elements in  $X$  (i.e.,  $H \in \mathbb{R}^{[0; \log(|X|)]}$ ; Cover and Thomas, 2006: 15 ???). For example, a *Bernoulli* ( $\mathcal{B}$ ) variable  $X$  with a probability of success of 0.5:  $X \sim \mathcal{B}(p = 0.5)$ , has an entropy of 1 bit. This is an upper bound for Bernoulli variables, since it decreases if  $p$  either increases or decreases, becoming 0 if  $p \in \{0; 1\}$ , making  $X$  constant and non-stochastic, and an entropy of 0 thus reflects that we know the values of  $X$  with absolute certainty (Cover and Thomas, 2006: 14–15). This example also provides the intuitive notion that the more uniform a distribution, the higher its entropy (Delgado-Bonal and Marshak, 2019).

While this conceptualization of entropy possess numerous desirable mathematical properties for measuring information, for example, being insensitive to the actual values of  $X$  and depend solely upon the their probabilities (Cover and Thomas, 2006: 14), it is, however, limited to *discrete* random variables. Since most phenomena of interest within the social sciences are *continuous* — probably because researchers often model such phenomena using generalized linear models (Field, 2018; Kruschke, 2014), which assume linearity (and thus continuity) in the parameters of the model (Wooldridge, 2019: 763) — it is necessary to extend this concept accordingly.

The *differential entropy* is ‘the entropy of a continuous random variable.’ (Cover and Thomas, 2006: 243) If  $X$  denotes a continuous random variable whose distribution can be characterized by the probability density function  $f(x)$ , the differential entropy,  $h(X)$ , is calculated with the formula (Ahmad and Lin, 1975: 373; Beirlant et al., 2001: 1; Berrett et al., 2019: 288; Cover and Thomas, 2006: 243, equation 8.1; Shannon, 1948):

$$h(X) \equiv - \int_S f(x) \log f(x) \, dx \quad (9)$$

where  $\int_S$  denote the integration operator applied over the *support set* of  $X$  with respect to  $x$  (i.e., where  $f(x) > 0$ ).<sup>10</sup> Unlike entropy, the differential entropy can be negative (Cover and Thomas, 2006: 244). An example of a continuous probability distribution is the normal (i.e.,  $\mathcal{N}(\mu, \sigma) \equiv \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$ ), whose differential entropy is easily calculated solely from its dispersion through the closed-form expression (Cover and Thomas, 2006: 244):

$$h(X) = \frac{1}{2} \log 2\pi e \sigma^2 \quad (10)$$

which for a standard normal variable becomes:  $\frac{1}{2} \log 2\pi e \approx 2.047$  bits.

<sup>9</sup>Popular alternatives to the base 2 for the logarithmic function are the base of the exponential  $e$  and 10, which means that the natural logarithm (i.e.,  $\ln$ ) is being used for the former. Instead of bits, this produces entropies on the scale of *natural units* (i.e., *nats*) and *decimal units* (i.e., *bats*), respectively (Back et al., 2018; Cover and Thomas, 2006: 14; Delgado-Bonal and Marshak, 2019).

<sup>10</sup>Since (9) involves an integral, the calculation of differential entropy presupposes that the density function and integral for  $X$  exists (Cover and Thomas, 2006: 243–244).

The concept of entropy can be found useful in relation to the states of knowledge previously conceptualized using Bayesian inference. Since that statistical framework similarly relies on the probability distributions of random variables to characterize prior and posterior knowledge, their concepts can be combined, leading to the notion that the different states of knowledge implied by the prior and posterior distributions simply express different states of information. Specifically, the prior distribution has an entropy, which express the degree of uncertainty in one's information about the parameter of interest  $\theta$ . The posterior distribution, by comparison, differs from the prior to the extent that the analyzed data has increased or reduced the uncertainty about one's information about  $\theta$ . For example, the normal prior distributions with a dispersion ( $\sigma$ ) of 1 previously specified as examples in section (3.1.1) have a differential entropy of approximately 2.05 bits, while those with a dispersion of 0.5 have a differential entropy of approximately 1.05 bits.

The example demonstrates that the concept of (differential) entropy can be imported to Bayesian inference and used to characterize the information contained within a state of knowledge, but it also makes it clear that a mere comparison of (differential) entropy between the prior and posterior would fail to consider changes in the mean of the distribution, since the absolute discrepancy between the differential entropy of the posterior previously mentioned in section (3.1.2) and the first and third prior (i.e.,  $|\Delta h| \approx 1$  bits), and the second and fourth prior (i.e.,  $|\Delta h| \approx 0$  bits), are the same, despite it being established that intuition would suggest that the degrees of discrepancy between them, and thus the surprise, should differ.

While the concept of (differential) entropy has thus shown to be a useful foundation to link concepts from Bayesian inference and Information theory, it provides an insufficient to measure surprise. This can be mitigated by introducing the final concept of relative entropy.

### 3.2.2 Relative entropy

The *relative entropy* quantifies information divergence by being 'a measure of the distance between two distributions.' (Cover and Thomas, 2006: 19) It is perhaps better known as the *Kullback-Leibler divergence* (KL divergence; e.g., Pérez-Cruz, 2019: 1), in reference to its original authors Kullback and Leibler (1951). Unlike entropy, which quantifies the uncertainty of a distribution, the relative entropy quantifies the divergence between distributions, with 'divergence' simply referring to the 'statistical distance' them (which should *not* be mistaken for a geometric distance; Nielsen, 2019: 4). It has proven 'central to information theory and statistics' (Pérez-Cruz, 2019: 1084; see also Cover and Thomas, 2006: 347-408) and seen application across numerous disciplines (e.g., Burnham and Anderson, 2001; Chaloner and Verdinelli, 1995).

For two discrete random variables  $p$  and  $q$ , the relative entropy is denoted  $D(p||q)$  and mathematically defined by formula (11) (Cover and Thomas, 2006: 19; Delgado-Bonal and Marshak, 2019: equation 5); and if  $p$  and  $q$  are continuous random variables with probability density functions  $f$  and  $g$ , it is given by formula (12) (Cover and Thomas, 2006: 251; Pérez-Cruz, 2019: 1):

$$D(p||q) \equiv \sum_{x \in X} p(x) \log \frac{p(x)}{q(x)} \quad (11)$$

$$\equiv \int f \log \frac{f}{g} \quad (12)$$

where for (11), it is taken by convention that  $0 \log 0 \equiv 0$ ,  $0 \log 0 \equiv 0$ , and  $p \log \infty \equiv \infty$ , which similarly applies to (12) (Cover

and Thomas, 2006: 19). While the properties of relative entropy for the discrete and continuous case are mostly the same (Cover and Thomas, 2006: 251), for example, they satisfy the property of *information monotonicity* by being invariant (???) (cf. Nielsen, 2019: 3), the relative entropy is zero if and only if (iff)  $p = q$  for the discrete case, and  $f = g$  (almost) everywhere for the continuous case (Cover and Thomas, 2006: 252-253; see also Pérez-Cruz, 2019: 1). For both instances, the relative entropy can also be infinite, with finite values only occurring for the continuous case when the support set of  $f$  is contained in the support set of  $g$  (i.e., if  $f$  is absolutely continuous with respect to  $g$ ). Thus, the relative entropy is generally a non-negative scalar on the extended real number line (i.e.,  $D(p||q) \in \bar{R}^{[0;\infty]}$ ).

While the relative entropy was developed to quantify how easy it is to discriminate between two distributions (Kullback and Leibler, 1951: 79), it is generally interpreted as a quantification of how inefficient it is to assume that a random variable is distributed as  $q$  when it is actually distributed as  $p$  (Cover and Thomas, 2006: 19). This has gained widespread appeal, since researchers may interpret the relative entropy as the mean of the *log-likelihood ratio* (Cover and Thomas, 2006: 375-380; Pérez-Cruz, 2019: 1), which makes it useful for hypothesis testing and model selection (see also, e.g., King, 1998; McElreath, 2019), since it can be used to justify preferring a statistical model with the lowest relative entropy on the basis of *parsimony*.

That interpretation can be recognized as similarly useful for the current purposes of conceptualizing surprise, since it may be meaningfully rephrased so that the relative entropy 'reflects the gain of information resulting from adding the knowledge in  $p$  relative to  $q$ , or the gain of information when we learn that the real probability function was  $p$  instead of  $q$  as we thought initially.' (Delgado-Bonal and Marshak, 2019) Such an interpretation fits the ongoing Bayesian (re)conceptualization of surprise by enabling the relative entropy between the prior and posterior distribution to be meaningfully interpreted as the gain in information resulting from adding the knowledge of the data to the prior. This means that divergence, by merit of being conceptually based in Bayesian inference and Information theory, coherently reflects how unexpected the results are compared to one's expectations, with zero divergence occurring iff the results exactly match one's expectations.

This leads to the completed, formal conceptualization of surprise as *the relative entropy between prior and posterior knowledge*. Within this conceptualization, knowledge and information are equated, while the prior and posterior are taken to reflect two states of knowledge, with the discrepancy between them being the degree of surprise. Higher degrees of discrepancy thus reflect greater surprise, measured in units of the *base* used for the logarithmic function, which here is specified as 2 so that the units of surprise are in bits.

This conceptualization fits perfectly with the informal notion of surprise identified among social science researchers (Section 2.1). This is evident, because the informal conceptualization of surprise involved a discrepancy between results and expectations/predictions about those results based in theory, which is directly reflected in this formal conceptualization, since the relative entropy measures the divergence (i.e., discrepancy) between the posterior knowledge (i.e., results) and the prior knowledge (i.e., predictions based in theory). It also reflects the *relative* and *dynamic* characteristics of surprise, since the degree of surprise within this formal conceptualization is always relative to the prior, with different degrees of surprise being possible if the posterior is compared to different prior specifications representing different theoretical expectations, and as the prior is continuously updated based on new data to better reflect reality, the relative entropy, and thus the surprise, will lessen. This formal (re)conceptualization of surprise thus offers a theoretically coherent understanding of surprise that fits with social

science researchers' intuitive use of the word, making it particularly suited for that field.

Now, perhaps attributable to its coherent conceptualization and intuitive interpretation, this is not the first time that relative entropy has been proposed as part of formals conceptualization of surprise across different scientific fields. For example, it has previously been proposed within the disciplines of clinical care (e.g., Burkhardt et al., 2025), ecology (e.g., Burnham and Anderson, 2001), software engineering (e.g., Bencomo and Belaggoun, 2014), and computer science (e.g., Itti and Baldi, 2009). However, while these proposals are undoubtedly helpful in reducing the mispractice of using the  $p$ -value as an implicit measure of surprise, their choice to not just conceptualize surprise as the relative entropy, but also as its measure, can be considered suboptimal when considering some of its limitations.

A limitation of relative entropy is that it is *not* a 'proper' metric in the geometric sence of the word. This is because it is *asymmetric*, which generally means that  $D(p||q) \neq D(q||p)$ , and it fails to satisfy the *triangle inequality* (Cover and Thomas, 2006: 19; Delgado-Bonal and Marshak, 2019), which means it lacks a series of intuitive properties that one may generally attribute to a distance. For example, suppose  $A$ ,  $B$ , and  $C$  represent points in Euclidean space, with the measure of distance travelled being asymmetric and in violation of the triangle inequality. Then, the distance travelled directly from  $A$  to  $C$  may not be the same as the distance travelled directly from  $C$  to  $A$ , and either of these may be longer than the distance travelled from  $A$  to  $B$  to  $C$ . This is counterintuitive (cf. Cover and Thomas, 2006: 367), since from real-world experience, one would expect distances to be symmetric, with stops and detours having no influence or adding to the distance travelled, respectively. These limitations thus make the relative entropy harder to interpret and invites researcher discretion by being sensitive to the arguably arbitrary choice of whether to calculate it between  $p$  and  $q$  or between  $q$  and  $p$ .

The relative entropy is by no means the only measure of divergence, with numerous alternatives (e.g., J divergence; Jeffreys, 1946), some of which are proper geometric measures of distance (Kullback and Leibler, 1951: 79). For example, the Bhattacharyya (1943; 1946) distances, the Earth mover's distance (Kantorovich, 1939; Monge, 1781; Vaserštejn, 1969), and the Mahalanobis' (1936) generalized distance. While each of these measures of distance possess their advantages and disadvantages, the one proposed here to measure surprise within the social sciences is the Jensen-Shannon divergence, that, while based on the relative entropy, possess none of its flaws.

Formally introduced by Lin (1991; see also Wong and You, 1985), the *Jensen-Shannon divergence* (JS divergence) is a popular alternative measure of divergence (Cha, 2007). Based on the KL divergence, it is mathematically defined with the formula (Lin, 1991: 146; Nielsen, 2019: 3):

$$\text{JS}(p||q) \equiv \frac{\text{D}(p||M) + \text{D}(q||M)}{2} \quad (13)$$

where  $M = \frac{p+q}{2}$  is the mixture distribution of  $p$  and  $q$ . From this formula, it can be recognized that the JS divergence is related to the relative entropy by being the mean of the relative entropies between  $p$  and  $q$  and their mixture distribution (Lin, 1991: 146). JS divergence shares numerous properties of the relative entropy useful to conceptualizing and measuring information divergence (Lin, 1991; Nielsen, 2019: 3), for example, being non-negative and zero iff  $p = q$  (Lin, 1991: 146-147).

While the JS divergence is closely related to the relative entropy, it possess numerous advantages. For ex-

ample, the JS divergence is always finite due to using the mixture distribution, since this means that  $p$  and  $q$  need not have overlapping support sets (i.e., they need *not* be absolutely continuous with respect to each other; Lin, 1991: 146-147). Another advantage is that the JS divergence generally is a non-negative, real-valued scalar with an upper-bound of  $\log 2$  (i.e.,  $\text{JS} \in \mathbb{R}^{[0:\log 2]}$ ) (Lin, 1991: 147; Nielsen, 2019: 3), which means that the upper bound is 1 if using base 2 for the logarithm. These bounds can be considered common for measures within the social sciences, for example, explained variance ( $R^2$ , Gelman et al., 2019), congeneric reliability ( $\rho_c$ ; Jöreskog, 1971), Sklar's  $\omega$  (Hughes, 2023), making this measure more intuitively interpretable within that discipline. Perhaps more importantly, unlike relative entropy, the JS divergence is symmetric, and by taking its square-root (i.e.,  $\sqrt{\text{JS}}$ ), one obtains the Jensen-Shannon distance (JS distance), which satisfies the triangle inequality (Nielsen, 2019: 3), making it a proper geometric measure of distance.

The JS distance can thus serve this formal conceptualization of surprise by providing the measurement of the degree of surprise between the prior and posterior knowledge. If the prior and posterior distribution are normal, then the JS distance is relatively easily computed using numerical integration. For example, if one again considers the normal prior distributions in figure (1) in section (3.1.1) with the posterior being equivalent to the fourth prior, then the discrepancy between the posterior and each prior can be computed using numerical integration as being approximately .684, .913, .910, and .000, respectively. As previously discussed, these values fit the conceptual expectations about surprise, since the discrepancy between the posterior and the fourth prior (Prior #4) is effectively zero, consistent with the complete absence of surprise when results match the theoretical expectations; the degree of surprise is higher when the theoretical expectation is more certain, as revealed by the higher discrepancy between the posterior and the narrow second prior (Prior #2) as compared to the wider first prior (Prior #1); and the degree of surprise is also higher if the theoretical expectation incorrectly predicts the magnitude and direction of the results, with the posterior-prior discrepancy being higher for the third prior (Prior #3) than the first prior (Prior #1), since the former expects the parameter of interest to be positive when it instead is negative. While this example presupposes that one knows the distribution of the prior and posterior, which in practice is generally not the case, this issue is relatively easily solved with the MCMC procedure of Bayesian inference, which further corroborates the utility of that framework as a basis for the conceptualization of surprise. By contrast, if the prior and posterior are not normal, however, it can be necessary to rely on density estimators (further discussed in section 4).

Substituting the relative entropy for the JS distance as the measure of surprise within this paper's formal conceptualization thus solves the potential issues of surprise otherwise possibly being infinite and difficult to interpret. That this is conceptually appropriate is further corroborated by the notion that the JS divergence was originally conceptualized in relation to Bayesian decision-making (Lin, 1991: 147-148), making it a natural choice for use in relation to a concept based on this framework.

## 4 Demonstration

[Work-in-progress]

## 5 Discussion

[Work-in-progress]

## 6 Conclusion

[Work-in-progress]

## 7 Acknowledgements

[Work-in-progress]

## References

- Abramowitz, A. I. and Webster, S. W. (2018). Negative partisanship: Why Americans dislike parties but behave like rabid partisans. *Political Psychology*, 39:119–135.
- Agresti, A. (2018). *Statistical Methods for the Social Sciences*. Pearson, 5 edition. Global Edition.
- Ahmad, I. A. and Lin, P.-E. (1975). A nonparametric estimation of the entropy for absolutely continuous distributions.
- Alford, J. R., Hatemi, P. K., Hibbing, J. R., Martin, N. G., and Eaves, L. J. (2011). The politics of mate choice. *Journal of Politics*, 73:362–379.
- Allaire, J. J., Francois, R., Ushey, K., Vandenbrouck, G., Geelnard, M., and Intel (2025). *RcppParallel: Parallel Programming Tools for 'Rcpp'*. R package version 5.1.10.
- Allaire, J. J., Xie, Y., Dervieux, C., McPherson, J., Luraschi, J., Ushey, K., Atkins, A., Wickham, H., Cheng, J., Chang, W., and Iannone, R. (2024). *rmarkdown: Dynamic Documents for R*. R package version 2.29.
- Altay, S., Hacquin, A.-S., and Mercier, H. (2020). Why do so few people share fake news? it hurts their reputation. *New Media and Society*.
- American Statistical Association (2016). American statistical association releases statement on statistical significance and p-values: Provides principles to improve the conduct and interpretation of quantitative science. Accessed March 18th, 2025.
- Auguie, B. (2017). *gridExtra: Miscellaneous Functions for "Grid" Graphics*. R package version 2.3.
- Bache, S. M. and Wickham, H. (2022). *magrittr: A Forward-Pipe Operator for R*. R package version 2.0.3.
- Back, A., Angus, D., and Wiles, J. (2018). Determining the number of samples required to estimate entropy in natural sequences. *arXiv*, 1805.08929v1:1–6.
- Barrett, T., Dowle, M., Srinivasan, A., Gorecki, J., Chirico, M., Hocking, T., Schwendinger, B., and Krylov, I. (2025). *data.table: Extension of 'data.frame'*. R package version 1.17.8.
- Bates, D., Mächler, M., Bolker, B., and Walker, S. (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, 67(1):1–48.
- Bates, D., Maechler, M., and Jagan, M. (2025). *Matrix: Sparse and Dense Matrix Classes and Methods*. R package version 1.7-2.

Bates, D., Mullen, K. M., Nash, J. C., and Varadhan, R. (2024). *minqa: Derivative-Free Optimization Algorithms by Quadratic Approximation*. R package version 1.2.8.

Bayes, T. and Price, R. (1763). An essay towards solving a problem in the doctrine of chance. By the late Rev. Mr. Bayes, communicated by Mr. Price, in a letter to John Canton, A.M.F.R.S. *Philosophical Transactions of the Royal Society of London*, 53:370–417.

Beirlant, J., Dudewicz, E. J., Györfi, L., and van der Meulen, E. C. (2001). Nonparametric entropy estimation: an overview. Unpublished manuscript.

Beleites, C. (2020). *arrayhelpers: Convenience Functions for Arrays*. R package version 1.1-0.

Ben-Shachar, M. S., Lüdecke, D., and Makowski, D. (2020). *effectsize*: Estimation of effect size indices and standardized parameters. *Journal of Open Source Software*, 5(56):2815.

Bencomo, N. and Belaggoun, A. (2014). A world full of surprises: Bayesian theory of surprise to quantify degrees of uncertainty. ICSE Companion 2014, pages 460–463, New York, NY, USA. Association for Computing Machinery.

Bengtsson, H. (2025). *matrixStats: Functions that Apply to Rows and Columns of Matrices (and to Vectors)*. R package version 1.5.0.

Berrett, T. B., Samworth, R. J., and Yuan, M. (2019). Efficient multivariate entropy estimation via  $k$ -nearest neighbour distances. *The Annals of Statistics*, 47(1):288–318.

Beygelzimer, A., Kakadet, S., Langford, J., Arya, S., Mount, D., and Li, S. (2024). *FNN: Fast Nearest Neighbor Search Algorithms and Applications*. R package version 1.1.4.1.

Bhattacharyya, A. K. (1943). On a measure of divergence between two statistical populations defined by their probability distributions. *Bulletin of the Calcutta Mathematical Society*, 35:99–109.

Bhattacharyya, A. K. (1946). On a measure of divergence between two multinomial populations. *Sankya*, 7(4):401–406.

Bishop, D. V. M. (2006). The psychology of experimental psychologists: Overcoming cognitive constraints to improve research: The 47th Sir Frederic Bartlett lecture. *Quarterly Journal of Experimental Psychology*, 73(1):1–19.

Bolker, B. (2025). *reformulas: Machinery for Processing Random Effect Formulas*. R package version 0.4.1.

Boshnakov, G. N. (2025). *Rdpack: Update and Manipulate Rd Documentation Objects*. R package version 2.6.4.

Boshnakov, G. N. and Putman, C. (2024). *rbibutils: Read 'Bibtex' Files and Convert Between Bibliography Formats*. R package version 2.3.

Brooks, S., Gelman, A., Jones, G., and Meng, X.-L. (2011). *Handbook of Markov Chain Monte Carlo*. Chapman and Hall, 1 edition.

- Burkhart, M. C., Ramadan, B., Solo, L., Parker, W. F., and Beaulieu-Jones, B. K. (2025). Quantifying surprise in clinical care: Detecting highly informative events in electronic health records with foundation models.
- Burnham, K. P. and Anderson, D. R. (2001). Kullback–leibler information as a basis for strong inference in ecological studies. *Wildlife Research*, 28(2):111–119.
- Buunk, A. P. and van Vugt, M. (2013). *Applying Social Psychology. From Problems to Solutions*. Sage, 2 edition.
- Bürkner, P.-C. (2017). brms: An R package for Bayesian multilevel models using Stan. *Journal of Statistical Software*, 80(1):1–28.
- Bürkner, P.-C. (2018). Advanced Bayesian multilevel modeling with the R package brms. *The R Journal*, 10(1):395–411.
- Bürkner, P.-C., Gabry, J., Kay, M., and Vehtari, A. (2025). posterior: Tools for working with posterior distributions. R package version 1.6.1.
- Cambridge Dictionary (2025). surprise.
- Canty, A. and Ripley, B. D. (2024). boot: Bootstrap R (S-Plus) Functions. R package version 1.3-31.
- Cha, S.-H. (2007). Comprehensive survey on distance/similarity measures between probability density functions. *International Journal of Mathematical Models and Methods in Applied Sciences*, 4(1):300–307.
- Chaloner, K. and Verdinelli, I. (1995). Bayesian experimental design: A review. *Statistical Science*, 10(3):273–304.
- Champely, S. (2020). pwr: Basic Functions for Power Analysis. R package version 1.3-0.
- Chang, W. (2024a). cachem: Cache R Objects with Automatic Pruning. R package version 1.1.0.
- Chang, W. (2024b). fastmap: Fast Data Structures. R package version 1.2.0.
- Chang, W. (2025). R6: Encapsulated Classes with Reference Semantics. R package version 2.6.1.
- Chang, W., Cheng, J., Allaire, J. J., Sievert, C., Schloerke, B., Xie, Y., Allen, J., McPherson, J., Dipert, A., and Borges, B. (2025a). shiny: Web Application Framework for R. R package version 1.11.1.
- Chang, W., Cheng, J., and Gao, C. (2025b). later: Utilities for Scheduling Functions to Execute Later with Event Loops. R package version 1.4.2.
- Cheng, J. (2025a). miniUI: Shiny UI Widgets for Small Screens. R package version 0.1.2.
- Cheng, J. (2025b). promises: Abstractions for Promise-Based Asynchronous Programming. R package version 1.3.3.
- Cheng, J., Chang, W., Reid, S., Brown, J., Trower, B., and Peslyak, A. (2025). httpuv: HTTP and WebSocket Server Library. R package version 1.6.16.

- Cheng, J., Sievert, C., Schloerke, B., Chang, W., Xie, Y., and Allen, J. (2024). *htmltools: Tools for HTML*. R package version 0.5.8.1.
- Christensen, G., Freese, J., and Miguel, E. (2019). *Transparent and Reproducible Social Science Research: How to do Open Science*. University of California Press, Oakland, California.
- Clayton, A. (2021). *Bernoulli's Fallacy: Statistical Illogic and the Crisis of Modern Science*. Columbia University.
- Cochrane, R., Billig, M., and Hogg, M. (1979). Politics and values in Britain: A test of Rokeach's two-value model. *British Journal of Clinical Psychology*, 18:159–167.
- Cohen, J. (1994). The Earth is round ( $p < .05$ ). *The American Psychologist*, 49(21):997–1003.
- Cole, S. R., Edwards, J. K., and Greenland, S. (2021). Surprise! *American Journal of Epidemiology*, 190(2):191–193.
- Coombes, K. R., Brock, G., Abrams, Z. B., and Abruzzo, L. V. (2019). Polychrome: Creating and assessing qualitative palettes with many colors. *Journal of Statistical Software, Code Snippets*, 90(1):1–23.
- Cover, T. M. and Thomas, J. A. (2006). *Elements of Information Theory*. John Wiley and Sons, 2 edition.
- Csárdi, G. (2019). *pkgconfig: Private Configuration for R Packages*. R package version 2.0.3.
- Csárdi, G. (2025). *cli: Helpers for Developing Command Line Interfaces*. R package version 3.6.5.
- Csárdi, G., Hester, J., Wickham, H., Chang, W., Morgan, M., and Tenenbaum, D. (2024). *remotes: R Package Installation from Remote Repositories, Including 'GitHub'*. R package version 2.5.0.
- Dahl, D. B., Scott, D., Roosen, C., Magnusson, A., and Swinton, J. (2019). *xtable: Export Tables to LaTeX or HTML*. R package version 1.8-4.
- Davison, A. C. and Hinkley, D. V. (1997). *Bootstrap Methods and Their Applications*. Cambridge University Press, Cambridge. ISBN 0-521-57391-2.
- Delgado-Bonal, A. and Marshak, A. (2019). Approximate entropy and sample entropy: A comprehensive tutorial. *Entropy*, 21:541.
- Dellsén, F. (2024). *Abductive Reasoning in Science*. Cambridge University Press.
- Eddelbuettel, D. (2013). *Seamless R and C++ Integration with Rcpp*. Springer, New York. ISBN 978-1-4614-6867-7.
- Eddelbuettel, D. (2017). *random: True Random Numbers using RANDOM.ORG*. R package version 0.2.6.
- Eddelbuettel, D. (2024). *digest: Create Compact Hash Digests of R Objects*. R package version 0.6.37.
- Eddelbuettel, D. and Balamuta, J. J. (2018). Extending R with C++: A Brief Introduction to Rcpp. *The American Statistician*, 72(1):28–36.

- Eddelbuettel, D. and François, R. (2011). Rcpp: Seamless R and C++ integration. *Journal of Statistical Software*, 40(8):1–18.
- Eddelbuettel, D., François, R., Allaire, J. J., Ushey, K., Kou, Q., Russell, N., Ucar, I., Bates, D., and Chambers, J. (2025). *Rcpp: Seamless R and C++ Integration*. R package version 1.1.0.
- Field, A. (2018). *Discovering Statistics Using IBM SPSS Statistics*. SAGE, 5 edition.
- Flake, J. K. and Fried, E. (2020). Measurement schmeasurement: Questionable measurement practices and how to avoid them. *Advances in Methods and Practices in Psychological Science*, 3(4):456–465.
- French, S. and Murphy, A. (2023). The value of surprise in science. *Erkenntnis*, 88:1447–1466.
- Friedman, J., Hastie, T., and Tibshirani, R. (2010). Regularization paths for generalized linear models via coordinate descent. *Journal of Statistical Software*, 33(1):1–22.
- Gabry, J., Goodrich, B., Lysy, M., and Johnson, A. (2024). *rstantools: Tools for Developing R Packages Interfacing with 'Stan'*. R package version 2.4.0.
- Gabry, J., Simpson, D., Vehtari, A., Betancourt, M., and Gelman, A. (2019). Visualization in Bayesian workflow. *The Journal of the Royal Statistical Society, Series A (Statistics in Society)*, 182:389–402.
- Gagolewski, M. (2022). stringi: Fast and portable character string processing in R. *Journal of Statistical Software*, 103(2):1–59.
- Gelman, A. (2024). A chain as strong as its strongest link? understanding the causes and consequences of biases arising from selective analysis and reporting of research results. *Journal of Research on Educational Effectiveness*, 17(3):459–461.
- Gelman, A., Carlin, J. B., Stern, H. S., Dunson, D. B., Vehtari, A., and Rubin, D. B. (2014). *Bayesian Data Analysis*. Chapman and Hall, 3 edition.
- Gelman, A., Goodrich, B., Gabry, J., and Vehtari, A. (2019). R-squared for Bayesian regression models. *The American Statistician*, 73(3):307–309.
- Gelman, A. and Stern, H. (2006). The difference between “significant” and “not significant” is not itself statistically significant. *The American Statistician*, 60(4):328–331.
- Genz, A. and Bretz, F. (2009). *Computation of Multivariate Normal and t Probabilities*. Lecture Notes in Statistics. Springer-Verlag, Heidelberg.
- Gigerenzer, G. (2004). Mindless statistics. *The Journal of Socio-Economics*, 33(5):587–606.
- Gigerenzer, G. (2018). Statistical rituals: The replication delusion and how we got there. *Advances in Methods and Practices in Psychological Science*, 1(2):198–218.
- Gordon, M., Gragg, S., and Konings, P. (2024). *htmlTable: Advanced Tables for Markdown/HTML*. R package version 2.4.3.

- Goren, P. and Chapp, C. (2017). Moral power: How public opinion on culture war issues shapes partisan predispositions and religious orientations. *The American Political Science Review*, 111(1):110–128.
- Greenland, S., Mansournia, M. A., and Joffe, M. (2022). To curb research misreporting, replace significance and confidence by compatibility. *Preventive Medicine*, 164:107127.
- Grolemund, G. and Wickham, H. (2011). Dates and times made easy with lubridate. *Journal of Statistical Software*, 40(3):1–25.
- Gronau, Q. F., Singmann, H., and Wagenmakers, E.-J. (2020). bridgesampling: An R package for estimating normalizing constants. *Journal of Statistical Software*, 92(10):1–29.
- Grosjean, P. (2025). *SciViews-R*. UMONS, MONS, Belgium.
- Gruber, J. (2014). *The Markdown File Extension*. The Daring Fireball Company, LLC.
- Grund, S., Robitzsch, A., and Luedtke, O. (2023). *mitml: Tools for Multiple Imputation in Multilevel Modeling*. R package version 0.4-5.
- Hankin, R. K. S. (2007). Very large numbers in R: Introducing package Brobdingnag. *R News*, 7.
- Harrell Jr, F. E. (2025). *Hmisc: Harrell Miscellaneous*. R package version 5.2-3.
- Hartley, R. V. L. (1928). Transmission of information. *The Bell System Technical Journal*, page 535.
- Hastie, T. J., Tibshirani, R. J., and Friedman, J. (2017). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer, 2 edition.
- Henry, L. and Wickham, H. (2023). *lifecycle: Manage the Life Cycle of your Package Functions*. R package version 1.0.4.
- Henry, L. and Wickham, H. (2024). *tidyselect: Select from a Set of Strings*. R package version 1.2.1.
- Henry, L. and Wickham, H. (2025). *rlang: Functions for Base Types and Core R and 'Tidyverse' Features*. R package version 1.1.6.
- Hester, J. and Bryan, J. (2024). *glue: Interpreted String Literals*. R package version 1.8.0.
- Hester, J., Henry, L., Müller, K., Ushey, K., Wickham, H., and Chang, W. (2024). *withr: Run Code 'With' Temporarily Modified Global State*. R package version 3.0.2.
- Hester, J., Wickham, H., and Csárdi, G. (2025). *fs: Cross-Platform File System Operations Based on 'libuv'*. R package version 1.6.6.
- Hughes, J. (2023). *sklarsomega: Measuring Agreement Using Sklar's Omega Coefficient*. R package version 3.0-2.
- Husserl, E. (1983). *Ideas Pertaining to a Pure Phenomenology and to a Phenomenological Philosophy*. Springer. First book.

- Ioannidis, J. P. A., Stanley, T. D., and Doucouliagos, H. (2017). The power of bias in economics research. *The Economic Journal (London)*, 127(605):F236–F265.
- Itti, L. and Baldi, P. (2009). Bayesian surprise attracts human attention. *Vision Research (Oxford)*, 49(10):1295–1306.
- Ivanova, M. and Vaidyanathan, B. (2024). Surprise in science: A qualitative study. *Erkenntnis*.
- Iverson, G. J., Lee, M. D., and Wagenmakers, E.-J. (2009). P(rep) misestimates the probability of replication. *Psychonomic Bulletin and Review*, 16(2):424.
- Jaynes, E. T. (1957a). Information theory and statistical mechanics. i. *The Physical Review*, 106(4):620–630.
- Jaynes, E. T. (1957b). Information theory and statistical mechanics. ii. *The Physical Review*, 108(2):171–190.
- Jaynes, E. T. (1968). Prior probabilities. *IEEE Transactions on Systems Science and Cybernetics*, ssc-4(3):227–241.
- Jeffreys, H. (1946). An invariant form for the prior probability in estimation problems. *Proceedings of the Royal Society, Series A*, 186:453–461.
- Johnson, A. R. (2025). *QuickJSR: Interface for the 'QuickJS-NG' Lightweight 'JavaScript' Engine*. R package version 1.8.0.
- Johnson, S. G. (2008). *The NLOpt nonlinear-optimization package*.
- Jung, L. (2024). *scrutiny: Error Detection in Science*. R package version 0.5.0.
- Jöreskog, K. G. (1971). Statistical analysis of sets of congeneric tests. *Psychometrika*, 36(2):109–133.
- Kant, I. (2018[1781]). *The Critique of Pure Reason*. CreateSpace Independent Publishing Platform. English translation by J. M. D. Meiklejohn.
- Kantorovich, L. V. (1939). Mathematical methods of organizing and planning production. *Management Science*, 6(4):366–422.
- Kay, M. (2024a). ggdist: Visualizations of distributions and uncertainty in the grammar of graphics. *IEEE Transactions on Visualization and Computer Graphics*, 30(1):414–424.
- Kay, M. (2024b). tidybayes: *Tidy Data and Geoms for Bayesian Models*. R package version 3.0.7.
- Kay, M. (2025). ggdist: *Visualizations of Distributions and Uncertainty*. R package version 3.3.3.
- Killeen, P. R. (2005a). An alternative to null-hypothesis significance tests. *Psychological Science*, 16(5):345–353.
- Killeen, P. R. (2005b). Replicability, confidence, and priors. *Psychological Science*, 16:1009–1012.
- King, G. (1998). *Unifying Political Methodology - The Likelihood Theory of Statistical Inference*. The University Of Michigan Press.

- Kruschke, J. K. (2014). *Doing Bayesian data analysis: A tutorial with R, JAGS, and Stan*. Academic Press, 2 edition.
- Kullback, S. and Leibler, R. A. (1951). On information and sufficiency. *The Annals of Mathematical Statistics*, 22(1):79–86.
- Lakatos, I. (1970). Falsification and the methodology of scientific research. In Lakatos, I. and Musgrave, A., editors, *Criticism and the Growth of Knowledge*, pages 91–196. Cambridge University Press.
- Lamport, L. (1986). *LATEX: a document preparation system*. Addison-Wesley Pub. Co.
- Lang, M. (2017). checkmate: Fast argument checks for defensive R programming. *The R Journal*, 9(1):437–445.
- Lang, M., Murdoch, D., and R Core Team (2024). *backports: Reimplementations of Functions Introduced Since R-3.0.0*. R package version 1.5.0.
- Laplace, P. S. (1951). *A Philosophical Essay on Probabilities*. Dover Publications, 6 edition. English translation by F. W. Truscott and F. L. Emory.
- Lenz, G. and Sahn, A. (2021). Achieving statistical significance with covariates and without transparency. *Political Analysis*, 29(3):356–369.
- Levy, R. and Mislevy, R. J. (2020). *Bayesian Psychometric Modeling*. Chapman and Hall.
- Ligges, U. and Mächler, M. (2003). Scatterplot3d - an R package for visualizing multivariate data. *Journal of Statistical Software*, 8(11):1–20.
- Lin, J. (1991). Divergence measures based on the Shannon entropy. *IEEE Transactions on Information Theory*, 37(1):145 – 151.
- Lugg, A. (1979). Laudan and the problem-solving approach to scientific progress and rationality. *Phil. Soc. Sci*, 9:466–474.
- Macdonald, R. R. (2005). Why replication probabilities depend on prior probability distributions. *Psychological Science*, 16:1006–1008.
- Maechler, M., Rousseeuw, P., Struyf, A., Hubert, M., and Hornik, K. (2024). *cluster: Cluster Analysis Basics and Extensions*. R package version 2.1.8 — For new features, see the 'NEWS' and the 'Changelog' file in the package source).
- Mahalanobis, P. C. (2018[1936]). On the generalized distance in statistics. *Sankhya, Series A*, 80:S1–S7.
- Makowski, D., Ben-Shachar, M. S., and Lüdecke, D. (2019). bayestestR: Describing effects and their uncertainty, existence and significance within the Bayesian framework. *Journal of Open Source Software*, 4(40):1541.
- Maraun, M. and Gabriel, S. (2010). Killeen's (2005) p rep coefficient: logical and mathematical problems. *Psychological Methods*, 15(2):182–191.

- Masi, M. (2023). Quantum indeterminism, free will, and self-causation. *Journal of Consciousness Studies*, 30(5):32–56.
- Mayo, D. G. (1996). *Error and the Growth of Experimental Knowledge*. The University of Chicago Press.
- McElreath, R. (2019). *Statistical Rethinking: A Bayesian Course with Examples in R*. Chapman and Hall, 2 edition.
- McShane, B. B., Gal, D., Gelman, A., Robert, C., and Tackett, J. L. (2019). Abandon statistical significance. *The American Statistician*, 73(sup1):235–245.
- Meehl, P. E. (1978). Theoretical risks and tabular asterisks: Sir Karl, Sir Ronald, and the slow progress of soft psychology. *Journal of Consulting and Clinical Psychology*, 46:806–834.
- Mello, P. R., Quartin, M., Schaefer, B. M., and Schosser, B. (2025). On the full non-Gaussian surprise statistic and the cosmological concordance between DESI, SDSS and Pantheon. *The Open Journal of Astrophysics*, 8.
- Microsoft (2024). *Windows 11*. x64, build 26100.
- Microsoft and Weston, S. (2022). *foreach: Provides Foreach Looping Construct*. R package version 1.5.2.
- Miller, D. (1974). Popper's qualitative theory of verisimilitude. *The British Journal for the Philosophy of Science*, 25(2):166–177.
- Modirshanechi, A., Brea, J., and Gerstner, W. (2021). Surprise: a unified theory and experimental predictions.
- Monge, G. (1781). *Mémoire sur la théorie des déblais et des remblais*. Imprimerie Royale.
- Morgan, M. S. (2005). Experiments versus models: New phenomena, inference and surprise. *Journal of Economic Methodology*, 12(2):317–329.
- Müller, K. (2023). *hms: Pretty Time of Day*. R package version 1.1.3.
- Müller, K. and Wickham, H. (2025a). *pillar: Coloured Formatting for Columns*. R package version 1.11.0.
- Müller, K. and Wickham, H. (2025b). *tibble: Simple Data Frames*. R package version 3.3.0.
- Neuwirth, E. (2022). *RColorBrewer: ColorBrewer Palettes*. R package version 1.1-3.
- Nielsen, F. (2019). On the Jensen–Shannon symmetrization of distances relying on abstract means. *entropy*, 21:485.
- Nyquist, H. (1924). Certain factors affecting telegraph speed. *The Bell System Technical Journal*, page 324.
- Nyquist, H. (1928). Certain topics in telegraph transmission theory. *A.I.E.E. Trans.*, 47:617.
- O'Hara-Wild, M., Kay, M., Hayes, A., and Hyndman, R. (2024). *distributional: Vectorised Probability Distributions*. R package version 0.5.0.
- Ooms, J. (2014). The jsonlite package: A practical and consistent mapping between JSON data and R objects. *arXiv:1403.2805 [stat.CO]*.

- Ooms, J. (2025a). *curl: A Modern and Flexible Web Client for R*. R package version 6.4.0.
- Ooms, J. (2025b). *V8: Embedded JavaScript and WebAssembly Engine for R*. R package version 7.0.0.
- Osmundsen, M., Bor, A., Vahlstrup, P. B., Bechmann, A., and Petersen, M. B. (2021). Partisan polarization is the primary psychological motivation behind political fake news sharing on twitter. *The American Political Science Review*, 115(3):999–1015.
- Pasek, J. (2025). *weights: Weighting and Weighted Statistics*. R package version 1.1.2.
- Patil, I., Makowski, D., Ben-Shachar, M. S., Wiernik, B. M., Bacher, E., and Lüdecke, D. (2022). datawizard: An R package for easy data preparation and statistical transformations. *Journal of Open Source Software*, 7(78):4684.
- Pedersen, T. L., Nicolae, B., and François, R. (2024). *farver: High Performance Colour Space Manipulation*. R package version 2.1.2.
- Peng, R. D. (2006). Interacting with data using the filehash package. *R News*, 6(4):19–24.
- Pennycook, G., McPhetres, J., Zhang, Y., Lu, J. G., and Rand, D. G. (2020). Fighting COVID-19 misinformation on social media: Experimental evidence for a scalable accuracy-nudge intervention. *Psychological Science*, 31(7):770–780.
- Pennycook, G. and Rand, D. G. (2019). Lazy, not biased: Susceptibility to partisan fake news is better explained by lack of reasoning than by motivated reasoning. *Cognition*, 188:39–50.
- Petersen, M. B. and Laustsen, L. (2019). Upper-body strength and political egalitarianism: Twelve conceptual replications. *Political Psychology*, 40(2):375–394.
- Pinheiro, J., Bates, D., and R Core Team (2025). *nlme: Linear and Nonlinear Mixed Effects Models*. R package version 3.1-167.
- Plate, T. and Heiberger, R. (2024). *abind: Combine Multidimensional Arrays*. R package version 1.4-8.
- Plummer, M., Best, N., Cowles, K., and Vines, K. (2006). Coda: Convergence diagnosis and output analysis for MCMC. *R News*, 6(1):7–11.
- Popper, K. (2002[1959]). *The Logic of Scientific Discovery*. Routledge, New York.
- Posit Team (2025). *RStudio: Integrated Development Environment for R*. Posit Software, PBC, Boston, MA.
- Pérez-Cruz, F. (2019). Kullback-Leibler divergence estimation of continuous distributions. *IEEE International Symposium on Information Theory*.
- Quartagno, M. and Carpenter, J. (2023). *jomo: A package for Multilevel Joint Modelling Multiple Imputation*.
- R Core Team (2025). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria.
- R Core team, Hester, J., and Csárdi, G. (2021). *urlchecker: Run CRAN URL Checks from Older R Versions*. R package version 1.0.1.

- Rafi, Z. and Greenland, S. (2020). Semantic and cognitive tools to aid statistical science: replace confidence and significance by compatibility and surprise. *BMC Medical Research Methodology*, 20(1):244.
- Revolution Analytics and Weston, S. (2022). *iterators: Provides Iterator Construct*. R package version 1.0.14.
- Rinker, T. W. and Kurkiewicz, D. (2018). *pacman: Package Management for R*. Buffalo, New York. version 0.5.0.
- Ritchie, S. (2020). *Science Fictions: How Fraud, Bias, Negligence, and Hype Undermine the Search for Truth*. Metropolitan Books.
- Robinson, D., Hayes, A., and Couch, S. (2025). *broom: Convert Statistical Objects into Tidy Tibbles*. R package version 1.0.8.
- Rokeach, M. (1973). *The nature of human values*. The Free Press.
- Rosenthal, R. (1979). File drawer problem and tolerance for null results. *Psychological Bulletin*, 86(3):638–641.
- Rényi, A. (1961). On measures of entropy and information. *Berkeley Symposium on Mathematical Statistics and Probability*, 4.1:547–561.
- Sarkar, D. (2008). *Lattice: Multivariate Data Visualization with R*. Springer, New York.
- Seehars, S., Amara, A., Refregier, A., Paranjape, A., and Akaret, J. (2014). Information gains from cosmic microwave background experiments. *Physical Review D, Particles, fields, gravitation, and cosmology*, 90(2):023533.
- Serlin, R. C. (2010). Regarding p rep: comment prompted by Iverson, Wagenmakers, and Lee (2010); Lecoutre, Lecoutre, and Poitevineau (2010); and Maraun and Gabriel (2010). *Psychological Methods*, 15(2):203–208.
- Shannon, C. E. (1948). A mathematical theory of communication. *The Bell System Technical Journal*, 27(3):379–423, 623–656.
- Shannon, C. E. (1964[1949]). The mathematical theory of communication. In Shannon, C. E. and Weaver, W., editors, *The Mathematical Theory of Communication*, pages 29–125. University of Illinois Press, Urbana.
- Sharpsteen, C. and Bracken, C. (2023). *tikzDevice: R Graphics Output in LaTeX Format*. R package version 0.12.6.
- Shaughnessy, J. J., Zechmeister, E. B., and Zechmeister, J. S. (2015). *Research Methods in Psychology*. McGraw-Hill, 10 edition. International Edition.
- Simmons, J. P., Nelson, L. D., and Simonsohn, U. (2011). False-positive psychology: Undisclosed flexibility in data collection and analysis allows presenting anything as significant. *Psychological Science*, 22(11):1359–1366.
- Sklyar, O., Eddelbuettel, D., and Francois, R. (2025). *inline: Functions to Inline C, C++, Fortran Function Calls from R*. R package version 0.3.21.

- Soetaert, K. (2024). *shape: Functions for Plotting Graphical Shapes, Colors*. R package version 1.4.6.1.
- Spinu, V. (2024). *timechange: Efficient Manipulation of Date-Times*. R package version 0.3.0.
- Stan Development Team (2020). StanHeaders: Headers for the R interface to Stan. R package version 2.21.0-6.
- Stan Development Team (2025). RStan: the R interface to Stan. R package version 2.32.7.
- Sterling, J., Jost, J. T., and Hardin, C. D. (2019). Liberal and conservative representations of the good society: A (social) structural topic modeling approach. *SAGE Open*, pages 1–13.
- Stock, J. and Watson, M. W. (2019). *Introduction to Econometrics*. Pearson Education Limited, 4 edition. Global Edition.
- The LaTeX Project (2025). *The LaTeX Project*. The LaTeX Project.
- Therneau, T. and Atkinson, B. (2025). *rpart: Recursive Partitioning and Regression Trees*. R package version 4.1.24.
- Therneau, T. M. (2024). *A Package for Survival Analysis in R*. R package version 3.8-3.
- Thành, H. T. and Knuth, D. E. (2025). *MiKTeX*. MiKTeX-pdfTeX version 4.21 (MiKTeX version 25.4).
- Tierney, L. (2024). *codetools: Code Analysis Tools for R*. R package version 0.2-20.
- Toshkov, D. (2016). *Research Design in Political Science*. Palgrave.
- Tukey, J. W. (1991). The philosophy of multiple comparisons. *Statistical Science*, 6:100–116.
- Urbanek, S. (2015). *base64enc: Tools for base64 encoding*. R package version 0.1-3.
- Ushey, K., Allaire, J. J., Wickham, H., and Ritchie, G. (2024). *rstudioapi: Safely Access the RStudio API*. R package version 0.17.1.
- Vaidyanathan, R., Xie, Y., Allaire, J., Cheng, J., Sievert, C., and Russell, K. (2023). *htmlwidgets: HTML Widgets for R*. R package version 1.6.4.
- van Buuren, S. and Groothuis-Oudshoorn, K. (2011). mice: Multivariate imputation by chained equations in R. *Journal of Statistical Software*, 45(3):1–67.
- van den Boogaart, K. G. (2023). *tensorA: Advanced Tensor Arithmetic with Named Indices*. R package version 0.36.2.1.
- Vaserštejn, L. N. (1969). Markov processes over denumerable products of spaces, describing large systems of automata. *Problemy Peredaci Informacii*, 5(3):64–72.
- Vaughan, D. (2025). *tzdb: Time Zone Database Information*. R package version 0.5.0.
- Vehtari, A., Gabry, J., Magnusson, M., Yao, Y., Bürkner, P.-C., Paananen, T., and Gelman, A. (2024). loo: Efficient leave-one-out cross-validation and WAIC for Bayesian models. R package version 2.8.0.

- Venables, W. N. and Ripley, B. D. (2002). *Modern Applied Statistics with S*. Springer, New York, 4 edition. ISBN 0-387-95457-0.
- Wagenmakers, E.-J., Lodewyckx, T., Kuriyal, H., and Grasman, R. (2010). Bayesian hypothesis testing for psychologists: A tutorial on the Savage-Dickey method. *Cognitive Psychology*, 60(3):158–189.
- Wand, M. (2025). *KernSmooth: Functions for Kernel Smoothing Supporting Wand & Jones (1995)*. R package version 2.23-26.
- Wang, Z., Zhang, H., Chen, J., and Chen, H. (2024). An effective framework for measuring the novelty of scientific articles through integrated topic modeling and cloud model. *Journal of Informetrics*, 18(4):101587.
- Warnes, G. R., Bolker, B., Lumley, T., Magnusson, A., Venables, B., Rydon, G., and Moeller, S. (2023). *gtools: Various R Programming Tools*. R package version 3.9.5.
- Warnes, G. R., Gorjanc, G., Magnusson, A., Andronic, L., Rogers, J., MacQueen, D., and Korosec, A. (2024). *gdata: Various R Programming Tools for Data Manipulation*. R package version 3.0.1.
- Wickham, H. (2016). *ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag New York.
- Wickham, H. (2021). *ellipsis: Tools for Working with ...* R package version 0.3.2.
- Wickham, H. (2023a). *forcats: Tools for Working with Categorical Variables (Factors)*. R package version 1.0.0.
- Wickham, H. (2023b). *stringr: Simple, Consistent Wrappers for Common String Operations*. R package version 1.5.1.
- Wickham, H., Averick, M., Bryan, J., Chang, W., McGowan, L. D., François, R., Grolemund, G., Hayes, A., Henry, L., Hester, J., Kuhn, M., Pedersen, T. L., Miller, E., Bache, S. M., Müller, K., Ooms, J., Robinson, D., Seidel, D. P., Spinu, V., Takahashi, K., Vaughan, D., Wilke, C., Woo, K., and Yutani, H. (2019). Welcome to the tidyverse. *Journal of Open Source Software*, 4(43):1686.
- Wickham, H., Bryan, J., Barrett, M., and Teucher, A. (2024a). *usethis: Automate Package and Project Setup*. R package version 3.1.0.
- Wickham, H., Chang, W., Flight, R., Müller, K., and Hester, J. (2025a). *sessioninfo: R Session Information*. R package version 1.2.3.
- Wickham, H., Chang, W., Hester, J., and Henry, L. (2024b). *pkgload: Simulate Package Installation and Attach*. R package version 1.4.0.
- Wickham, H., Chang, W., Luraschi, J., and Mastny, T. (2024c). *profvis: Interactive Visualizations for Profiling R Code*. R package version 0.4.0.
- Wickham, H., François, R., Henry, L., Müller, K., and Vaughan, D. (2023a). *dplyr: A Grammar of Data Manipulation*. R package version 1.1.4.
- Wickham, H. and Henry, L. (2025). *purrr: Functional Programming Tools*. R package version 1.1.0.

- Wickham, H., Henry, L., and Vaughan, D. (2023b). *vctrs: Vector Helpers*. R package version 0.6.5.
- Wickham, H., Hester, J., and Bryan, J. (2024d). *readr: Read Rectangular Text Data*. R package version 2.1.5.
- Wickham, H., Hester, J., Chang, W., and Bryan, J. (2022). *devtools: Tools to Make Developing R Packages Easier*. R package version 2.4.5.
- Wickham, H., Hester, J., Chang, W., Müller, K., and Cook, D. (2021). *memoise: 'Memoisation' of Functions*. R package version 2.0.1.
- Wickham, H., Hester, J., and Csárdi, G. (2025b). *pkgrbuild: Find Tools Needed to Build R Packages*. R package version 1.4.8.
- Wickham, H., Kuhn, M., and Vaughan, D. (2025c). *generics: Common S3 Generics not Provided by Base R Methods Related to Model Fitting*. R package version 0.1.4.
- Wickham, H. and Pedersen, T. L. (2024). *gttable: Arrange 'Grobs' in Tables*. R package version 0.3.6.
- Wickham, H., Pedersen, T. L., and Seidel, D. (2025d). *scales: Scale Functions for Visualization*. R package version 1.4.0.
- Wickham, H., Vaughan, D., and Girlich, M. (2024e). *tidyverse: Tidy Messy Data*. R package version 1.3.1.
- Wickham, H. and Xie, Y. (2025). *evaluate: Parsing and Evaluation Tools that Provide More Details than the Default*. R package version 1.0.4.
- Wiener, N. (2019[1961]). *Cybernetics: Or Control and Communication in the Animal and the Machine*. The MIT Press, 2 edition. Reissue of the 1961 second edition.
- Wong, A. K. C. and You, M. (1985). Entropy and distance of random graphs with application to structural pattern recognition. *IEEE Trans. Pattern Anal. Machine Intell.*, PAMI-7(5):599–609.
- Wooldridge, J. M. (2019). *Introductory Econometrics*. Cengage, 7 edition.
- Xie, Y. (2014). knitr: A comprehensive tool for reproducible research in R. In Stodden, V., Leisch, F., and Peng, R. D., editors, *Implementing Reproducible Computational Research*. Chapman and Hall/CRC. ISBN 978-1466561595.
- Xie, Y. (2015). *Dynamic Documents with R and knitr*. Chapman and Hall/CRC, Boca Raton, Florida, 2 edition. ISBN 978-1498716963.
- Xie, Y. (2016). *bookdown: Authoring Books and Technical Documents with R Markdown*. Chapman and Hall/CRC, Boca Raton, Florida.
- Xie, Y. (2025a). *bookdown: Authoring Books and Technical Documents with R Markdown*. R package version 0.43.
- Xie, Y. (2025b). *knitr: A General-Purpose Package for Dynamic Report Generation in R*. R package version 1.50.

- Xie, Y. (2025c). *mime: Map Filenames to MIME Types*. R package version 0.13.
- Xie, Y. (2025d). *xfun: Supporting Functions for Packages Maintained by 'Yihui Xie'*. R package version 0.52.
- Zeileis, A. and Croissant, Y. (2010). Extended model formulas in R: Multiple parts and multiple responses. *Journal of Statistical Software*, 34(1):1–13.
- Zeileis, A., Fisher, J. C., Hornik, K., Ihaka, R., McWhite, C. D., Murrell, P., Stauffer, R., and Wilke, C. O. (2020). *colorspace: A toolbox for manipulating and assessing colors and palettes*. *Journal of Statistical Software*, 96(1):1–49.
- Zeileis, A., Hornik, K., and Murrell, P. (2009). Escaping RGBland: Selecting colors for statistical graphics. *Computational Statistics & Data Analysis*, 53(9):3259–3270.
- Zhao, J. H. and Schafer, J. L. (2023). *pan: Multiple imputation for multivariate panel or clustered data*. R package version 1.9.

## 8 Appendix

This appendix contains the specifications for the software used in this paper (Section 8.1) and an R function implementing the KDE plug-in estimator that was used to compute the Jensen-Shannon distance from MCMC samples (Section 8.2) as demonstrated in section (4).

### 8.1 Software Specifications

This document was written with *markdown* (Gruber, 2014) and *LATeX*(Lamport, 1986; The LaTeX Project, 2025), using *pdfTeX*(v3.141592653-2.6-1.40.2; *MiKTeX*[v25.4], Thành and Knuth, 2025). Compiled on a Windows 11 x64 (build 26200) operating system (OS; Microsoft, 2024) on the 2026-01-14 15:45 CET (Timezone: Europe/Copenhagen), the machine has 16 cores available, and the range of doubles is  $2.23 \times 10^{-308}$  -  $1.8 \times 10^{+308}$ .

All data processing and analyses were conducted on the aforementioned software using the R programming language (v4.4.3 [2025-02-28 ucrt]; R Core Team, 2025) in the RStudio IDE (2025.5.1.513; Posit Team, 2025), which besides the base R packages by R Core Team (2025) used the following packages (including their dependencies, which were *not* necessarily used): *abind* (v1.4.8; Plate and Heiberger, 2024), *arrayhelpers* (v1.1.0; Beleites, 2020), *backports* (v1.5.0; Lang et al., 2024), *bayestestR* (v0.16.1; Makowski et al., 2019), *base64enc* (v0.1.3; Urbanek, 2015), *bayesplot* (v1.13.0; Gabry et al., 2019), *bookdown* (v0.43; Xie, 2016, 2025a), *boot* (v1.3.31; Canty and Ripley, 2024; Davison and Hinkley, 1997), *bridgesampling* (v1.1.2; Gronau et al., 2020), *brms* (v2.22.0; Bürkner, 2017, 2018), *Broddingnag* (v1.2.9; Hankin, 2007), *broom* (v1.0.8; Robinson et al., 2025), *cachem* (v1.1.0; Chang, 2024a), *checkmate* (v2.3.3; Lang, 2017), *cli* (v3.6.5; Csárdi, 2025), *cluster* (v2.1.8; Maechler et al., 2024), *coda* (v0.19.4.1; Plummer et al., 2006), *codetools* (v0.2.20; Tierney, 2024), *colorspace* (v2.1.1; Zeileis et al., 2009, 2020), *curl* (v6.4.0; Ooms, 2025a), *data.table* (v1.17.8; Barrett et al., 2025), *datawizard* (v1.1.0; Patil et al., 2022), *devtools* (v2.4.5; Wickham et al., 2022), *digest* (v0.6.37; Eddelbuettel, 2024), *distributional* (v0.5.0; O’Hara-Wild et al., 2024), *dplyr* (v1.1.4; Wickham et al., 2023a), *effectsize* (v1.0.1; Ben-Shachar et al., 2020), *ellipsis* (v0.3.2; Wickham, 2021), *evaluate* (v1.0.4; Wickham and Xie, 2025), *farver* (v2.1.2; Pedersen et al., 2024), *fastmap* (v1.2.0; Chang, 2024b), *filehash* (v2.4.6;

Peng, 2006), **FNN** (v1.1.4.1; Beygelzimer et al., 2024), **forcats** (v1.0.0; Wickham, 2023a), **foreach** (v1.5.2; Microsoft and Weston, 2022), **Formula** (v1.2.5; Zeileis and Croissant, 2010), **fs** (v1.6.6; Hester et al., 2025), **gdata** (v3.0.1; Warnes et al., 2024), **generics** (v0.1.4; Wickham et al., 2025c), **ggdist** (v3.3.3; Kay, 2024a, 2025), **ggplot2** (v3.5.2; Wickham, 2016), **glmnet** (v4.1.9; Friedman et al., 2010), **glue** (v1.8.0; Hester and Bryan, 2024), **gridExtra** (v2.3; Auguie, 2017), **gtable** (v0.3.6; Wickham and Pedersen, 2024), **gtools** (v3.9.5; Warnes et al., 2023), **Hmisc** (v5.2.3; Harrell Jr, 2025), **hms** (v1.1.3; Müller, 2023), **htmlTable** (v2.4.3; Gordon et al., 2024), **htmltools** (v0.5.8.1; Cheng et al., 2024), **htmlwidgets** (v1.6.4; Vaidyanathan et al., 2023), **httpuv** (v1.6.16; Cheng et al., 2025), **inline** (v0.3.21; Sklyar et al., 2025), **iterators** (v1.0.14; Revolution Analytics and Weston, 2022), **jomo** (v2.7.6; Quartagno and Carpenter, 2023), **jsonlite** (v2.0.0; Ooms, 2014), **KernSmooth** (v2.23.26; Wand, 2025), **knitr** (v1.50; Xie, 2014, 2015, 2025b), **later** (v1.4.2; Chang et al., 2025b), **lattice** (v0.22.6; Sarkar, 2008), **lifecycle** (v1.0.4; Henry and Wickham, 2023), **lme4** (v1.1.37; Bates et al., 2015), **loo** (v2.8.0; Vehtari et al., 2024), **lubridate** (v1.9.4; Grolemund and Wickham, 2011), **magrittr** (v2.0.3; Bache and Wickham, 2022), **MASS** (v7.3.64; Venables and Ripley, 2002), **Matrix** (v1.7.2; Bates et al., 2025), **matrixStats** (v1.5.0; Bengtsson, 2025), **memoise** (v2.0.1; Wickham et al., 2021), **mice** (v3.18.0; van Buuren and Groothuis-Oudshoorn, 2011), **mime** (v0.13; Xie, 2025c), **miniUI** (v0.1.2; Cheng, 2025a), **minqa** (v1.2.8; Bates et al., 2024), **mitml** (v0.4.5; Grund et al., 2023), **mvtnorm** (v1.3.3; Genz and Bretz, 2009), **nlme** (v3.1.167; Pinheiro et al., 2025), **nloptr** (v2.2.1; Johnson, 2008), **nnet** (v7.3.20; Venables and Ripley, 2002), **pacman** (v0.5.1; Rinker and Kurkiewicz, 2018), **pan** (v1.9; Zhao and Schafer, 2023), **pillar** (v1.11.0; Müller and Wickham, 2025a), **pkgbuild** (v1.4.8; Wickham et al., 2025b), **pkgconfig** (v2.0.3; Csárdi, 2019), **pkgload** (v1.4.0; Wickham et al., 2024b), **Polychrome** (v1.5.4; Coombes et al., 2019), **posterior** (v1.6.1; Bürkner et al., 2025), **profvis** (v0.4.0; Wickham et al., 2024c), **promises** (v1.3.3; Cheng, 2025b), **purrr** (v1.1.0; Wickham and Henry, 2025), **pwr** (v1.3.0; Champely, 2020), **QuickJSR** (v1.8.0; Johnson, 2025), **R6** (v2.6.1; Chang, 2025), **rbibutils** (v2.3; Boshnakov and Putman, 2024), **RColorBrewer** (v1.1.3; Neuwirth, 2022), **Rcpp** (v1.1.0; Eddelbuettel, 2013; Eddelbuettel and Balamuta, 2018; Eddelbuettel and François, 2011; Eddelbuettel et al., 2025), **RcppParallel** (v5.1.10; Allaire et al., 2025), **Rdpack** (v2.6.4; Boshnakov, 2025), **readr** (v2.1.5; Wickham et al., 2024d), **reformulas** (v0.4.1; Bolker, 2025), **remotes** (v2.5.0; Csárdi et al., 2024), **rlang** (v1.1.6; Henry and Wickham, 2025), **rmarkdown** (v2.29; Allaire et al., 2024), **rpart** (v4.1.24; Therneau and Atkinson, 2025), **rstan** (v2.32.7; Stan Development Team, 2025), **rstantools** (v2.4.0; Gabry et al., 2024), **rstudioapi** (v0.17.1; Ushey et al., 2024), **scales** (v1.4.0; Wickham et al., 2025d), **scatterplot3d** (v0.3.44; Ligges and Mächler, 2003), **scrutiny** (v0.5.0; Jung, 2024), **sessioninfo** (v1.2.3; Wickham et al., 2025a), **shape** (v1.4.6.1; Soetaert, 2024), **shiny** (v1.11.1; Chang et al., 2025a), **StanHeaders** (v2.32.10; Stan Development Team, 2020), **stringi** (v1.8.7; Gagolewski, 2022), **stringr** (v1.5.1; Wickham, 2023b), **survival** (v3.8.3; Therneau, 2024), **svUnit** (v1.0.6; Grosjean, 2025), **tensorA** (v0.36.2.1; van den Boogaart, 2023), **tibble** (v3.3.0; Müller and Wickham, 2025b), **tidybayes** (v3.0.7; Kay, 2024b), **tidyverse** (v2.0.0; Wickham et al., 2019), **tikzDevice** (v0.12.6; Sharpsteen and Bracken, 2023), **timechange** (v0.3.0; Spinu, 2024), **tzdb** (v0.5.0; Vaughan, 2025), **urlchecker** (v1.0.1; R Core team et al., 2021), **usethis** (v3.1.0; Wickham et al., 2024a), **v8** (v7.0.0; Ooms, 2025b), **vctrs** (v0.6.5; Wickham et al., 2023b), **weights** (v1.1.2; Pasek, 2025), **withr** (v3.0.2; Hester et al., 2024), **xfun** (v0.52; Xie, 2025d), and **xtable** (v1.8.4; Dahl et al., 2019). For the random number generator used by MCMC, the seed (i.e., 896663432) was exogenously

predetermined by the random R package (v0.2.6; Eddelbuettel, 2017).

## 8.2 Jensen-Shannon distance function

An implementation of the Jensen-Shannon distance function using R (v4.4.3 [2025-02-28 ucrt]; R Core Team, 2025), and an example of its application in relation to the posterior and prior from Bayesian inference in relation to empirical data using the brms R package (v2.22.0; Bürkner, 2017, 2018), is provided in the code below:

```

1 # Function for computing the Jensen-Shannon (JS) distance/divergence based on MCMC
  ↳ samples from the posterior and prior distributions using default parameter
  ↳ values
2 #
3 # @param p An n-dimensional vector of distributional samples
4 # @param q An n-dimensional vector of distributional samples
5 # @param n An integer indicating the length of the grid used for estimating the
  ↳ densities
6 # @param epsilon A small, positive real-valued scalar used to avoid instances of
  ↳ division by zero
7 # @param base A small, positive real-valued scalar used as the base of the
  ↳ logarithmic function
8 # @param type A character that determines whether the output is the JS distance or
  ↳ divergence
9 # @return A non-negative real-valued scalar reflecting the estimated discrepancy
  ↳ between the distributions
10
11 # install.packages("KernSmooth")
12 JSD <- function(p = NULL, q = NULL, n = 10000, epsilon = 1e-05, base = 2, type = c(
  ↳ "Distance", "Divergence")){
13   # Parameter checks
14   stopifnot(length(p)==length(q))
15   stopifnot(is.numeric(p))
16   stopifnot(is.numeric(q))
17   stopifnot(is.numeric(n))
18   stopifnot(n > 1)
19   stopifnot(is.numeric(epsilon))
20   stopifnot(epsilon > 0)
21   stopifnot(is.numeric(base))
22   stopifnot(base > 0)
23   type <- match.arg(type)
24
25   # Define a common grid
26   grid <- seq(min(c(p, q)), max(c(p, q)), length.out = n)
27
28   # Use Kernsmooth::bkde to estimate densities on the same grid if available,
```

```
    ↪ otherwise use stats::density
29 if(requireNamespace("KernSmooth", quietly = TRUE)){
30   require(KernSmooth)
31   p_density <- KernSmooth::bkde(p, gridsize = n, range.x = c(min(c(p, q)), max(c(
32     ↪ p, q))))$y
33   q_density <- KernSmooth::bkde(q, gridsize = n, range.x = c(min(c(p, q)), max(c(
34     ↪ p, q))))$y
35 }
36 else{
37   p_density <- stats::density(p, bw = "SJ", kernel = "gaussian", from = min(grid)
38     ↪ , to = max(grid), n = length(grid))$y
39   q_density <- stats::density(q, bw = "SJ", kernel = "gaussian", from = min(grid)
40     ↪ , to = max(grid), n = length(grid))$y
41 }
42
43 # Avoid any zeros
44 p_density <- pmax(p_density, epsilon)
45 q_density <- pmax(q_density, epsilon)
46
47 # Ensure normalized densities
48 dx <- diff(grid)[1]
49 p_density <- p_density / sum(p_density * dx)
50 q_density <- q_density / sum(q_density * dx)
51
52 # Compute mixture
53 m <- 0.5 * (p_density + q_density)
54
55 # Compute JSD
56 jensen_shannon_divergence <- 0.5 * sum(p_density * log(p_density / m) * dx) +
57   0.5 * sum(q_density * log(q_density / m) * dx)
58
59 # Change to user-specified base
60 jensen_shannon_divergence/log(base)
61
62 if(type == "Distance"){
63   jensen_shannon_distance <- sqrt(jensen_shannon_divergence)
64   return(jensen_shannon_distance)
65 }
66 else{
67   return(jensen_shannon_divergence)
68 }
```

```
1 # Example
2 # install.packages("tidyverse")
3 library("tidyverse")
4
5 # Load data
6 # install.packages("devtools")
7 # library("devtools")
8 # devtools::install_github("kosukeimai/qss-package", build_vignettes = TRUE)
9 data(face, package = "qss")
10 df <- face; rm(face)
11
12 # Recode data
13 df <- df %>%
14   mutate(
15     d.share = d.votes / (d.votes + r.votes),
16     r.share = r.votes / (d.votes + r.votes),
17     d.diff.share = d.share - r.share,
18     d.diff.share = standardize(d.diff.share),
19     d.comp = standardize(d.comp)
20   ) %>%
21   select(d.comp, d.diff.share) %>%
22   rename(
23     democratic_competence = d.comp,
24     democratic_relative_support = d.diff.share
25   )
26
27 # Standardize data matrix
28 # install.packages("datawizard")
29 library("datawizard")
30 df <- apply(df, 2, datawizard::standardize) %>% as.data.frame()
31
32 # Specify MCMC
33 # install.packages("parallel")
34 library("parallel")
35 posterior_samples <- 4000
36 chains <- cores <- parallel::detectCores()-1
37 warmup <- 1000
38 iter <- ceiling((posterior_samples / chains) + warmup)
39 posterior_samples <- (iter - warmup)*chains
40
41 # Fit BGML
42 # install.packages("brms")
```

```
43 library("brms")
44 democratic_relative_support_competence_fit <- brms::brm(
45   democratic_relative_support ~ 0 + democratic_competence,
46   data = df,
47   prior = c(
48     brms::set_prior("normal(0, 0.1)", class = "b", lb = -1, ub = 1),
49     brms::set_prior("normal(1, 0.1)", class = "sigma", lb = 0, ub = 1)
50   ),
51   family = gaussian(
52     link = "identity"
53   ),
54   sample_prior = "yes",
55   cores = cores,
56   chains = chains,
57   seed = 896663432,
58   iter = iter,
59   warmup = warmup
60 )
61 # summary(democratic_relative_support_competence_fit)
62
63 # Extract draws
64 democratic_relative_support_competence_draws <- brms::as_draws_df(
65   ↪ democratic_relative_support_competence_fit)
66 b_democratic_competence_posterior <-
67   ↪ democratic_relative_support_competence_draws$b_democratic_competence
68 b_democratic_competence_prior <-
69   ↪ democratic_relative_support_competence_draws$prior_b
70
71 # Compute JSKL divergence
72 JSD(b_democratic_competence_prior, b_democratic_competence_posterior)
```